

Red Hat OpenShift

OpenShift 4.x DevEx Basics Workshop

Jaroslaw Stakun

Version 1.0.0, March 05, 2020

Table of Contents

Introduction	1
Attendee details	1
What is Openshift	1
Links	1
Contributors	1
Application Basics	2
Introduction	2
Starting up - logging on and creating a project	2
Creating your first Application	2
Adding additional Applications	7
Interacting with OpenShift through the Command Line	8
Option 1: Installing the CLI on localhost	9
Option 2: Use a pre-configured docker image on OpenShift	9
Using your terminal (both options)	10
A Summary of Application Interactions	12
Understanding Deployments and Configuration	14
Introduction	14
Introducing Deployment Configurations	14
Dependency Injection using Config Maps	17
Dependency Injection of sensitive information using Secrets	21
Understanding the Deployment Strategies	22
Application Deployment Configurations and DevOps Approaches	24
Introduction	24
Workshop Content	24
Creation of version 1	24
Creation of version 2	26
Blue / Green Deployment	27
A/B Deployment	29
Cleaning up	31
Understanding the Software Defined Network within OpenShift	32
Introduction	32
The basics of Service Addressing	32
Using the shorthand Service name directly	37
Using the Fully Qualified Domain Name for accessing Services	38
Controlling Access through Network Policies	39
The RBAC model for Developers	42

Introduction.....	42
Examining Service Accounts	42
Adding Role Bindings to your namespace/project	45
Giving Users lower levels of permission	46
Understanding Persistent Volumes	48
Introduction.....	48
Adding a Persistent Volume to an Application	49
Demonstrating survivability of removal of all Pods	53
Pod Health Probes	54
Introduction.....	54
Creating the project and application.....	54
Viewing the running application	57
Liveness Probe	59
Activation of the Liveness Probe	62
Readiness Probe	63
Activation of the Readiness Probe	65
Cleaning up	67

Introduction

Attendee details

Name:	
User ID (userX):	

This workshop is designed to introduce Developers to **OpenShift 4** and explain the usage and technologies around it from a developer perspective.

What is Openshift

Red Hat® OpenShift® is a hybrid cloud, enterprise, secure Kubernetes application platform.

OpenShift is an Enterprise strength, secure implementation of the Kubernetes container orchestration project with additional tooling designed to make the lives of Developer and Administrators as easy as possible.

Links

- <https://www.openshift.com/learn/what-is-openshift>
- <https://en.wikipedia.org/wiki/OpenShift>
- <https://www.openshift.com/products/container-platform>
- <https://console-openshift-console.apps-crc.testing/> (the Web Console URL for the Workshop)
- <https://api.crc.testing:6443> (The Web Terminal URL for the workshop)

Contributors

Red Hat UK DSA Team:

- Phil Prosser
- Mark Roberts
- 'Uther' Lawson

Original asciidoc project for DevEx3.x by Phillip Kruger, Red Hat CEMEA

Application Basics

Introduction

This chapter will introduce the attendee to the mechanics of creating and manipulating Applications using OpenShift Container Platform. It will introduce the two distinct contexts of the User Interface, the Administrator and Developer views, and guide the attendee through creating some Applications.



The attendee will be using a browser for interacting with both the OpenShift Container Platform UI and the provided terminal (for command line exercises). It is suggested the attendee uses Chrome or Firefox.

Starting up - logging on and creating a project

Log on to cluster at <https://console-openshift-console.apps-crc.testing/> as *userx*, where x is the number assigned to you by the course administrator, password *openshift*

Ensure you are on the Administrator View - at the top left of the UI is a selection box which describes the current view selected.

The Administrator view provides you with an extended functionality interface that allows you to deep dive into the objects available to your user. The Developer view is an opinionated interface designed to ease the use of the system for developers. This workshop will have you swapping between the contexts for different tasks.

Click on *Create Project*

Name - ‘sandboxX’ where x is user number

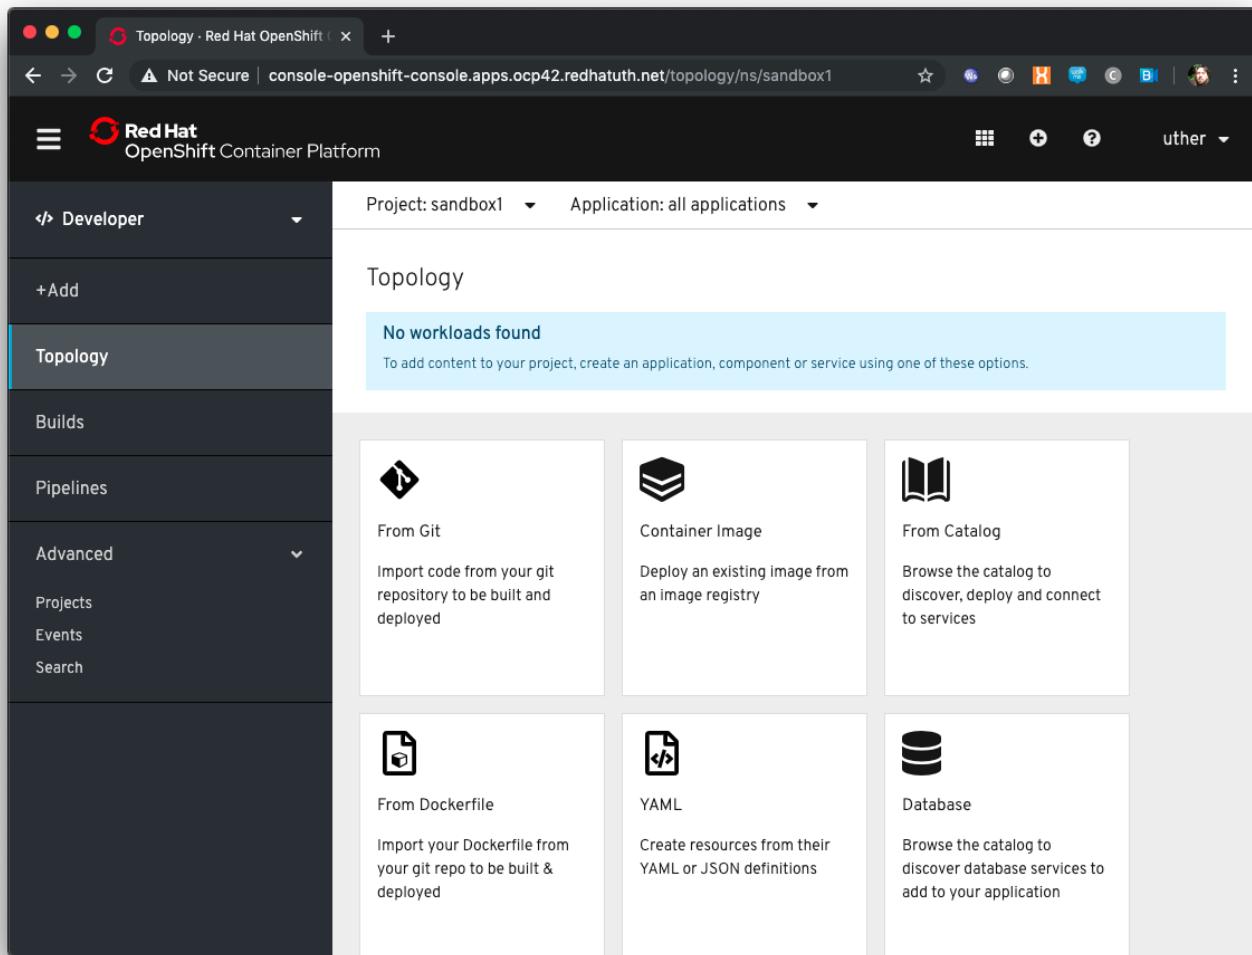
Display Name, Description are labels

When created click on *Role Bindings*

By default when you create a Project within OpenShift your user is given administration rights. This allows the user to create any objects that they have rights to create and to change the security and access settings for the project itself, i.e. add users as Administrators, Edit Access, Read access or disable other user’s abilities to even see the project and the objects within.

Creating your first Application

In the top left of the UI, where the label indicates the view mode, change the mode from Administrator to Developer



Click *Add*

The Catalog screen for the developer combines all the ways components can be added into the Project. These are:

- From Git - this provides another way to do a Source-2-Image build by first choosing the Git repo and then the builder image to use
- Container Image - this provides a way to directly deploy an Image from a repository
- From Catalog - this allows the Developer to browse all available templates, which are predefined sets of Objects to create an application
- From Dockerfile - this allows the Developer to do a controlled build of an Image from a Dockerfile
- YAML - this allows the Developer to provide a set of populated YAML to define the objects to be added to the Project
- Database - this allows the Developer to browse pre-created Database services to add to the Project

Select 'From Catalog'

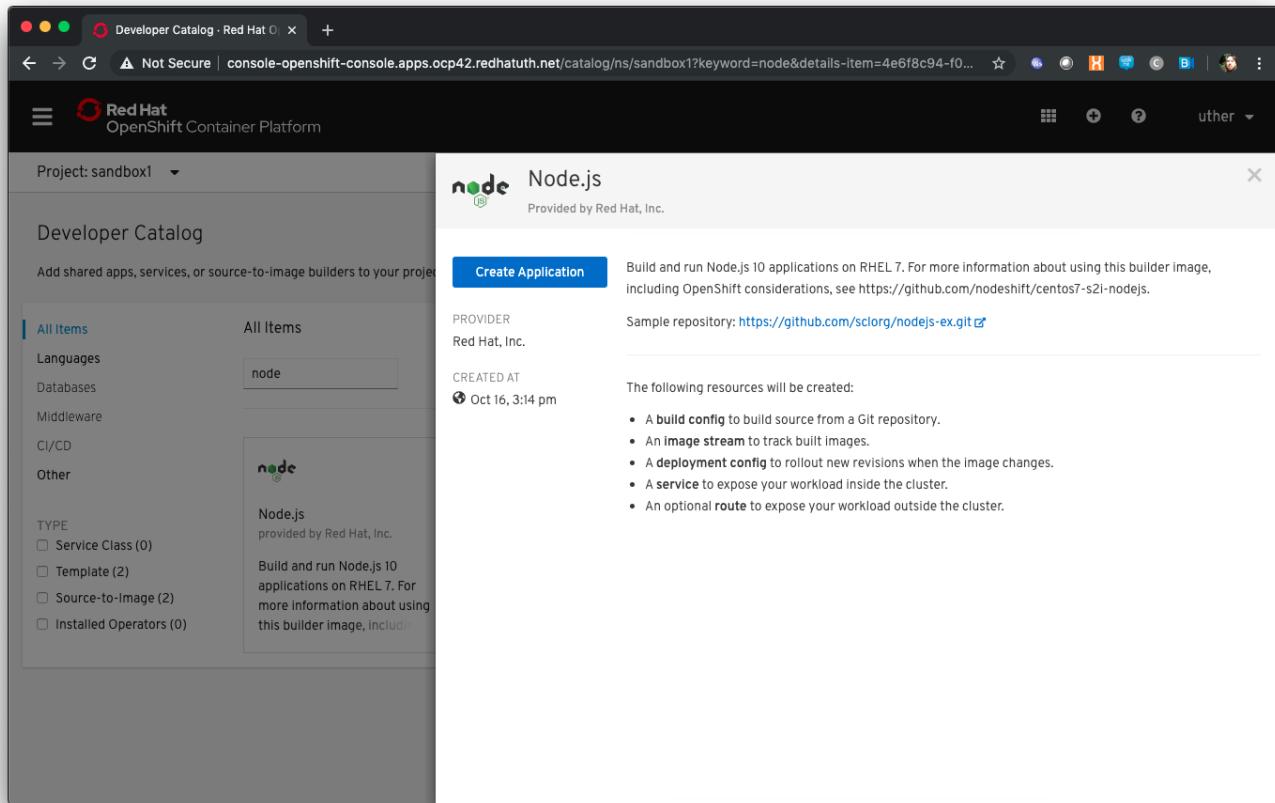
Enter 'node' in the search box

The screenshot shows the Red Hat OpenShift Container Platform Developer Catalog interface. The left sidebar has sections for Developer, +Add, Topology, Builds, Pipelines, and Advanced. Under Advanced, there is a TYPE section with Service Class (0), Template (2), Source-to-image (2), and Installed Operators (0). The main area is titled 'Developer Catalog' and says 'Add shared apps, services, or source-to-image builders to your project from the Developer Catalog. Cluster admins can install additional apps which will show up here automatically.' A search bar at the top right contains 'node'. Below it, a table lists four items:

Image	Name	Description
	Node.js	provided by Red Hat, Inc. Build and run Node.js 10 applications on RHEL 7. For more information about using this builder image, includ...
	Node.js + MongoDB	provided by Red Hat, Inc. An example Node.js application with a MongoDB database. For more information about using th...
	Node.js + MongoDB (Ephemeral)	provided by Red Hat, Inc. An example Node.js application with a MongoDB database. For more
	Tech Preview - Modern Web Applications	provided by Red Hat, Inc. Build and run Modern Web Applications on RHEL 7. For more information about us...

OpenShift allows for multiple base images to be built upon - the selection of node searches for any images or templates registered into the system with the label *node*. In the screenshot above, and in the catalog you will be presented with, there will be a selection of base images.

Select 'Node.js'



When you select an option, in this case the Node.js builder one, you are presented with a wizard that shows you exactly what components will be created as part of your Project. In this case, with Node.js, the template will create a build config, that will build the Image that will contain your Application, an ImageStream which is the OpenShift representation of an Image, a deployment config, which defines exactly how the image will be deployed as a running Container within the Project, a service which is the internal endpoint for the application within the Project and a route, optionally, which will provide access to the Application for external consumers.

Click on *Create Application*

This approach uses the OpenShift '*source-to-image*' concept, which is a controlled mechanism provided by OpenShift that automates the creation of application images using a base image and a source repository.

Change the Builder Image Version to 8

The '*source-to-image*' approach allows you to use differing versions of a base image - in this case you can execute the Node build against a number of supported Node versions.

Enter the following for the Git repo for the application - <https://github.com/utherp0/nodenews>

In a separate browser tab go to <https://github.com/utherp0/nodenews>

If you visit the URL you will see there is no OpenShift specific code in the repository at all.

Close the github tab

Back at the OCP4.2 UX leave the Application as ‘Create Application’.

OpenShift 4 introduces the concept of Application Grouping. This allows a visualisation of multiple Applications in the same group, making visibility of the Application much simpler.

Ensure the application name is ‘nodenews-app’

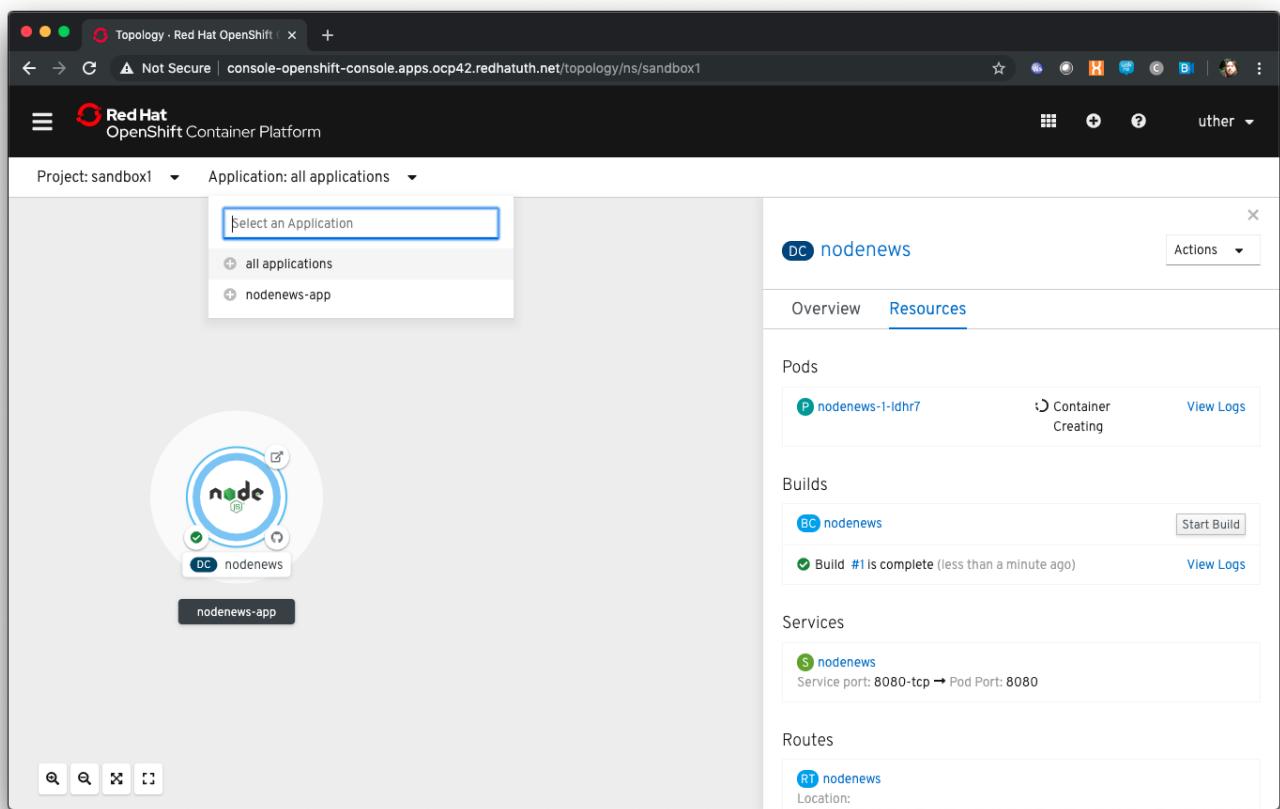
Ensure the Name is ‘nodenews’

Ensure the ‘Create Route is checked’

Click *Create*

The Topology view is a new feature of OpenShift 4 that provides a dynamic and useful visualisation of all of your Applications in a given Project.

Click on the Icon marked *Node*



The side-panel contains an overview of the Application you chose. In this case it will cover the build. Once a build has completed this side panel shows the Pods that are running, the builds that have completed, the services exported for the Application and the routes, if the Application has any.

Wait for the Build to finish, the Pod to change from Container Creating to Running

When an Application is created the Pod ring will be empty, indicating that an Application will appear once the build has completed. When the build completes the Pod ring will switch to light blue, indicating the Pod is being pulled (the image is being pulled from the registry to the Node where the Pod will land) and is starting (the Pod is physically in place but the Containers within it are not reporting as ready). Once the Pod is placed and running the colour of the Pod ring will change to dark blue.

Click on the Tick at the bottom left of the Pod

If you scroll the log of the Build output you will see the steps that the build takes. This includes laying the foundational file layers for the base image, performing the code specific build operations (in this case an '*npm install*') and then pushing the file layers for the image into the OpenShift integrated registry.

Adding additional Applications

Click on *Topology*

Click *Add*

Click *From Catalog*

Search for 'httpd'

Select the (httpd) template - if there is more than one you want to select the one that is a **builder** and not the example application one

Click on *Create Application*

Leave Image Version as 2.4

Set the git repo to '<https://github.com/utherp0/forumstaticassets>'

Make sure the Application is 'node-news-app'

Make sure the Name is forumstaticassets

Make sure the 'Create a Route' checkbox is clicked

Click *Create*

Note that the new Application icon appears within a bounded area on the Topology page labelled with the *Application* chosen above. If you click on the area between the Pods you can move the group as a single action.

Click on the forumstaticassets Pod

Watch the build complete, the Container Creating and the Running event

Click *Add*

Click *From Catalog*

Search for ‘node’

Select ‘Node.js’

Click *Create Application*

Leave at Builder Image Version 10

Set the git repo to ‘<https://github.com/utherp0/ocpnode>’

In the ‘Application’ pulldown select ‘Create Application’

In the ‘Application Name’ enter ‘ocpnode-app’

Ensure the Name is ‘ocpnode’

Ensure the ‘Create Route’ is checked

Click *Create*

Click on the ‘ocpnode’ Application in the topology - click on the ‘cross’ icon situated at the bottom left of the Topology panel to centralise the topology

Now we have created a new Application grouping you will see two ‘cloud’ groupings, labelled with the appropriate Application name you entered.

Interacting with OpenShift through the Command Line

With the OpenShift Enterprise command line interface (CLI), you can create applications and manage OpenShift projects from a terminal. The CLI is ideal in situations where you are:

- Working directly with project source code.
- Scripting OpenShift Enterprise operations.
- Restricted by bandwidth resources and cannot use the web console.

The CLI is available using the oc command:

```
oc {command}
```

For the duration of the course you will be using a provided Application hosted within OpenShift itself -

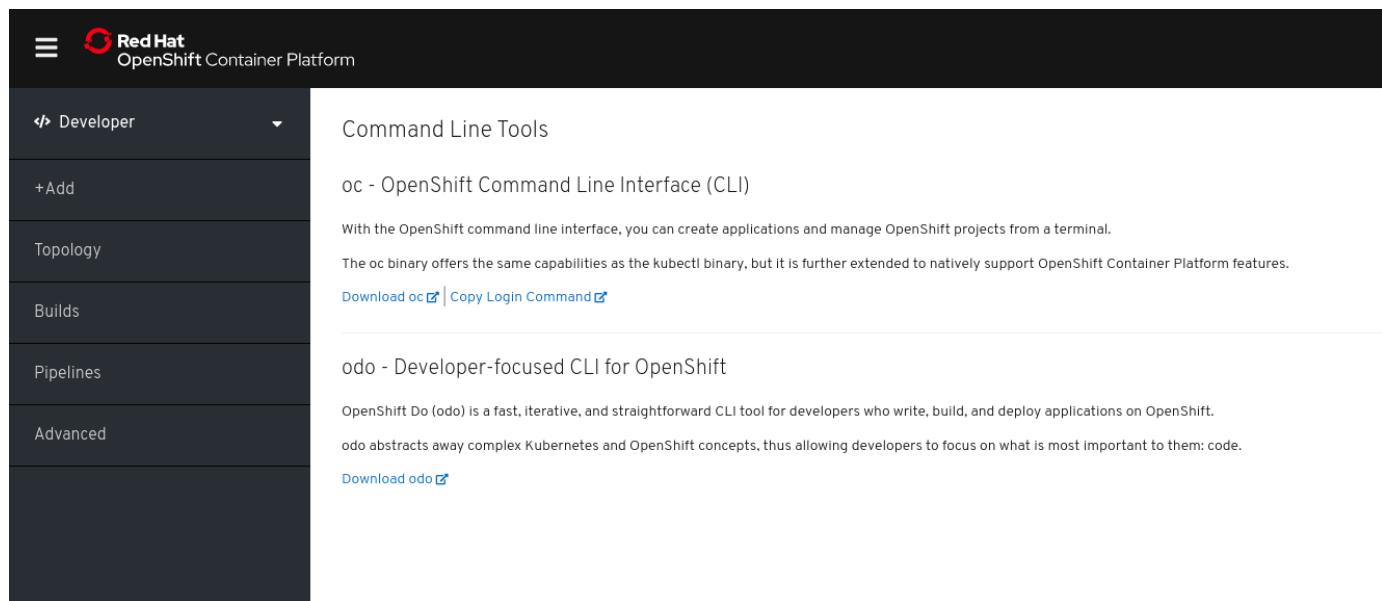
this is reachable at <https://api.crc.testing:6443>

The first time you visit the terminal application you will be prompted for a logon - use the same credentials as your OpenShift logon, userx and openshift

If you want to optionally install the oc client or use the terminal applications plus additional command line tools, such as ODO and tkn, you have two options; either install `oc` on your localhost or use a docker image in OpenShift. The image option is pre-built and running already and is easier if you have a locked-down laptop and are unable to download and install additional applications.

Option 1: Installing the CLI on localhost

The easiest way to download the CLI is by accessing the Command Line Tools page on the web console from the question mark on the top right corner [(?) → Command Line Tools]



The screenshot shows the Red Hat OpenShift Container Platform web interface. In the top left, there's a navigation bar with a menu icon, the Red Hat logo, and the text "Red Hat OpenShift Container Platform". On the left side, there's a sidebar with a dropdown menu set to "Developer", followed by links for "+Add", "Topology", "Builds", "Pipelines", and "Advanced". The main content area is titled "Command Line Tools". It contains two sections: "oc - OpenShift Command Line Interface (CLI)" and "odo - Developer-focused CLI for OpenShift". Each section includes a brief description, a "Download" link, and a "Copy Login Command" link. The "oc" section also has a note about its extended capabilities compared to kubectl.

From this page you can download the oc and odo command line tools.

If you wish to install the command line tools locally then you may also need to download and command line tools for OpenShift Pipelines (Tekton) and Kamel. Instructions on this is at the start of each chapter.

Option 2: Use a pre-configured docker image on OpenShift

This option installs and runs a docker image inside OpenShift that already has oc (and other command line tools) installed and configured. To install this image, do the following:

Login to the OCP system through the UI. (userX / openshift) Ensure that you are in the Developer mode of the user interface by selecting Developer in the top left corner menu. Select the project drop down menu and select Create Project.

Create a new project:

Name: oconline-userX

Display Name: OC Terminal

You will then see a Topology view that allows you to create new workloads from a variety of sources as shown below:

The screenshot shows the Red Hat OpenShift Container Platform Web Console. The top navigation bar includes the Red Hat logo and the text "Red Hat OpenShift Container Platform". The left sidebar has a "Developer" section with "Topology" selected, along with other options like "Builds", "Pipelines", and "Advanced". The main content area is titled "Topology" and displays a message "No workloads found" followed by the instruction "To add content to your project, create an application, component or service using one of these options." Below this are six cards representing different creation methods:

- From Git**: Import code from your git repository to be built and deployed.
- Container Image**: Deploy an existing image from an image registry.
- From Catalog**: Browse the catalog to discover, deploy and connect to services.
- From Dockerfile**: Import your Dockerfile from your git repo to be built & deployed.
- YAML**: Create resources from their YAML or JSON definitions.
- Database**: Browse the catalog to discover database services to add to your application.

Select container image and paste into the image name :

quay.io/marrober/devex-terminal-4

Press the search button on the end of the line to force the application to examine the repository image.

For the Application name enter : DevEx-Terminal

For the Name enter : devex-terminal

Ensure the box is ticked for *Create a route to the application*

Click create.

Wait for the pod to spin up and show a dark blue ring on the topology view.

Click on the pod and on then on the resources tab on the right hand side click on the route. This will launch the terminal in a new browser window.

Keep this terminal window open for the remainder of the workshop for all of your command line interaction.

Using your terminal (both options)

If this document refers to *your terminal* it will either be on your localhost or the containerised terminal depending on your option above.

(Login options for the installed OC command line tool only). * In the Web Console, select the top right

pulldown, choose **[Copy Login Command]** * This will require you to login again using your credentials (to prove who you are) and will then display a web page with *Display Token* on it. Select this link and copy the full link that begins with `oc login --token.....` * Paste this into the command line window on your local machine where you have access to run the `oc` command. Note that `oc` must be added to your path.

The response to executing the login command should look similar to that shown below :

```
Logged into "https://api.cluster-london-a6e1.london-a6e1.example.opentlc.com:6443" as  
"user1" using the token provided.
```

```
You don't have any projects. You can try to create a new project, by running
```

```
oc new-project <projectname>
```

Make sure `oc` is working, type:

```
oc whoami  
oc version
```



Also see the **Command-Line Reference**: https://docs.openshift.com/container-platform/4.2/cli_reference/openshift_cli/getting-started-cli.html

To explore the command line further execute the following commands and observe the results.

```
oc projects
```

User should have access to one project, explain the ability to access multiple projects

```
oc project sandboxX
```

User should now be using the `sandboxX` project created and configured earlier

Next we will try a command that will fail because of OpenShift's security controls

```
oc get users
```

There is a level of permission within the OpenShift system called '*Cluster Admin*'. This permission allows a User to access any of the objects on the system regardless of Project. It is effectively a super-user and as such normal users do not normally have this level of access.

```
oc get pods
```

If you look carefully at the Pods shown you will notice there are additional Pods above and beyond the ones expected for your Applications. If you look at the state of these Pods they will be marked as Completed. Everything in OpenShift is executed as a Pod, including Builds. These completed Pods are the Builds we have run so far.

```
oc get pods | grep Completed
```

```
oc get pods | grep Running
```

```
oc get dc
```

DC is an abbreviation for Deployment Config. These are Objects that directly define how an Application is deployed within OpenShift. This is the 'ops' side of the OpenShift system. Deployment Configs are different to Kubernetes Deployments in that they are an extension and contain things such as Config Maps, Secrets, Volume Mounts, labelled targetting of Nodes and the like.

Enter the command below to tell OpenShift to scale the number of instances of the Deployment Config *nodenews* to two rather than the default one.

```
oc scale dc/nodenews --replicas=2
```

A Summary of Application Interactions

Go back to the UI and make sure you are on Developer mode. Click on Topology.

Click on the 'nodenews' application

Note the 'DC' reference to the application under the icon

In the pop-up panel on the right click on *Resources*

Note that there are two pods running with the application now

Change the mode from Developer to Administrator

Select the *sandboxx* project in the project list

Note the metrics for the project

Click on *Workloads* and then select Pods.

Change to Developer mode and then select Topology if the Topology page isn't already shown

Click and hold on the forumstaticassets Pod and drag it onto the ocpnode-app

Select 'move' when it prompts whether you want to move it

This concept of Application grouping allows for much better visualisation of the complete application.

Continue on with the Deployments chapter, which uses the applications created here to show the capabilities of the deployment configuration and how to alter the behaviour and file system of a Container without changing the image.

Understanding Deployments and Configuration

Introduction

This chapter will introduce the attendee to concepts of deployments and configuration injection using OpenShift Container Platform.



The attendee will be using a browser for interacting with both the OpenShift Container Platform UI and the provided terminal (for command line exercises). It is suggested the attendee uses Chrome or Firefox.



This chapter follows directly on from the Application Basics one so it is assumed the attendee is logged on to the system and has the applications pre-created. If you are starting at this chapter please repeat the commands from the Application Basics chapter. You should start this chapter with two application running in the project, nodenews and ocpnode

Introducing Deployment Configurations

Ensure you are on the Administrator View (top level, select *Administrator*) Select *Workloads/Deployment Configs*

Name	Namespace	Labels	Status	Pod Selector
DC forumstaticassets	NS sandbox99	app=forumstaticassets app.kubernetes.io/c... =forumstatic... app.kubernetes.io/... =forumstatic... app.kubernetes.io/name=http app.kubernetes.io/par...=ocpnode-... app.openshift.io/runtime=http app.openshift.io/runtime-version=2.4	1 of 1 pods	Q app=forumstaticassets, deploymentconfig=forumstaticassets
DC nodenews	NS sandbox99	app=nodenews app.kubernetes.io/compo... =noden... app.kubernetes.io/insta... =nodenews... app.kubernetes.io/name=nodejs app.kubernetes.io/pa...=nodenews-... app.openshift.io/runtime=nodejs app.openshift.io/runtime-version=10	1 of 1 pods	Q app=nodenews, deploymentconfig=nodenews
DC ocpnode	NS sandbox99	app=ocpnode app.kubernetes.io/compo... =ocpn... app.kubernetes.io/instance=ocpnode app.kubernetes.io/name=nodejs app.kubernetes.io/par...=ocpnode-...	1 of 1 pods	Q app=ocpnode, deploymentconfig=ocpnode

Observe the information provided on the Deployment Configs screen. DC indicates the Deployment Config by name, NS is the project/namespace, Labels are the OpenShift labels applied to the DC for collating and visualisation, status indicates the active Pod count for the Deployment and pod selector indicates the labels used for the Pods attributed to the Deployment Config.

Click on the *nodenews* deployment-config

Next to the Pod icon, with the count in it which should be set to two, click on the up arrow twice

The screenshot shows the Red Hat OpenShift Container Platform web interface. The left sidebar is titled "Administrator" and includes sections for Home, Operators, Workloads (Pods, Deployments, Deployment Configs, Stateful Sets, Secrets, Config Maps), Cron Jobs, Jobs, Daemon Sets, Replica Sets, Replication Controllers, Horizontal Pod Autoscalers, Networking, and Storage. The "Deployment Configs" section is currently selected. The main content area is titled "Deployment Configs > Deployment Config Details" for "nodenews". It shows the "Overview" tab selected, with a summary: "Deployment Config Overview" showing a circular progress bar at 2, with an upward arrow indicating it's scaling to 4. Below this, there are detailed settings: Name (nodenews), Namespace (sandbox99), Labels (app=nodenews, app.kubernetes.io/component=nodenews, app.kubernetes.io/instance=nodenews, app.kubernetes.io/name=nodejs, app.kubernetes.io/part-of=nodenews-app, app.openshift.io/runtime=nodejs, app.openshift.io/runtime-version=10), Pod Selector (app=nodenews, deploymentconfig=nodenews), Latest Version (1), Reason (config change), Update Strategy (Rolling), Timeout (600 seconds), and Update Period (1 second). There are also tabs for "YAML", "Pods", "Environment", and "Events".

We have told the deployment config to run four replicas. The system now updated the replication controller, whose job is to make sure that x replicas are running healthily, and once the replication controller is updated the system will ensure that number of replicas is running

Click on 'YAML'

In the YAML find an entry for replicas. It should say 4. Set it to 2, then save.

```

31   resources: {}
32     activeDeadlineSeconds: 21600
33   triggers:
34     - type: ImageChange
35       imageChangeParams:
36         automatic: true
37         containerNames:
38           - nodenews
39         from:
40           kind: ImageStreamTag
41           namespace: sandbox99
42           name: nodenews:latest
43           lastTriggeredImage: >-
44             image-registry.openshift-image-registry.svc:5000/sandbox99/nodenews@sha256:9b49cd987fa9d7d3fd159512827808001716045fa9d405ebf3a3640d05b3e
45       - type: ConfigChange
46     replicas: 4
47     revisionHistoryLimit: 10
48     test: false
49     selector:
50       app: nodenews
51       deploymentconfig: nodenews
52     template:
53       metadata:
54         creationTimestamp: null
55         labels:

```

Save Reload Cancel Download

Click on ‘Deployment Configs’ near the top of the panel.

Select the nodenews deployment config by clicking on the nodenews DC link

Ensure the Pod count is now two

Dependency Injection using Config Maps

Click on *Workloads/Config Maps*

Config Maps allow you to create groups of name/value pairs in objects that can be expressed into the Applications via the Deployment Configs and can change the file system, environment and behaviour of the Application without having to change the *image* from whence the Application was created. This is extremely useful for dependency injection without having to rebuild the Application image

Click on ‘Create Config Map’

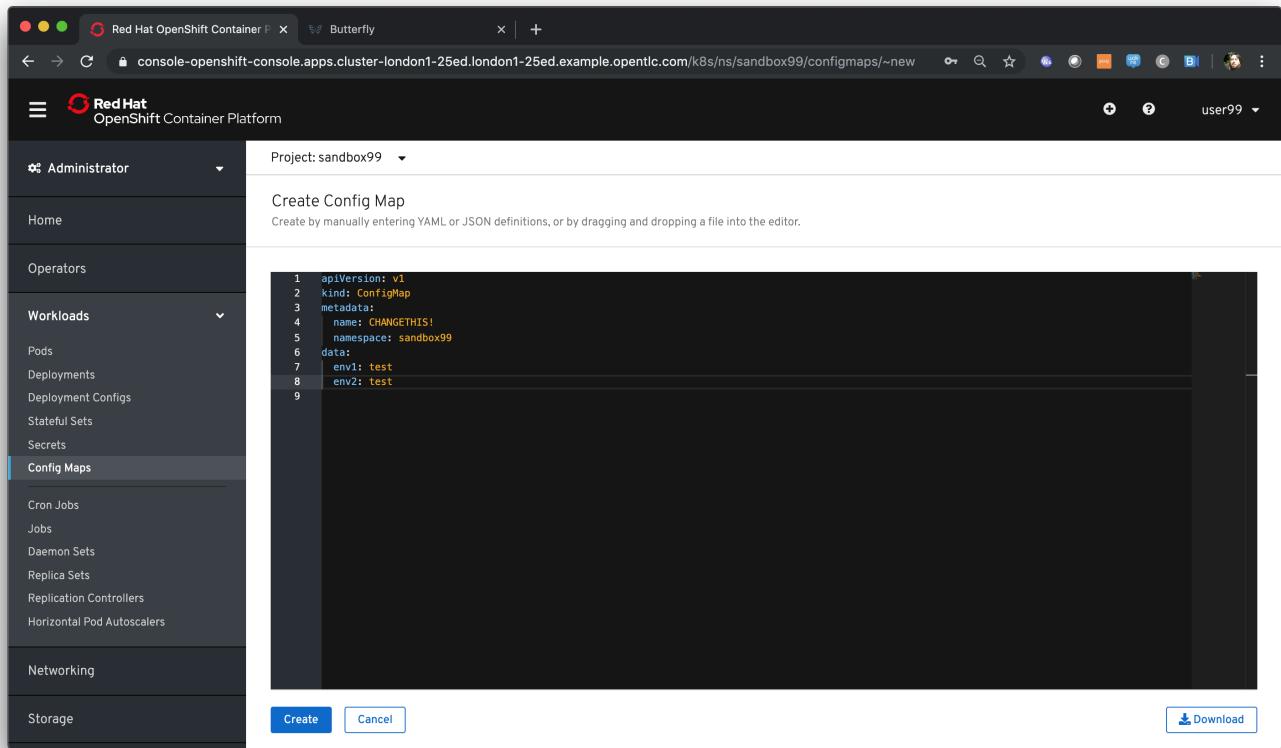
In the editor change the name: from example to ‘nodenewsconfig’

Delete all of the lines below **data:**

Add “ env1: test”

Add “ env2: test”

The YAML should look similar to the screenshot below, with your project name instead of sandbox99 and the name you have entered rather than CHANGETHIS.



```
1 apiVersion: v1
2 kind: ConfigMap
3 metadata:
4   name: CHANGETHIS!
5   namespace: sandbox99
6 data:
7   env1: test
8   env2: test
9
```

The 'Create' button is visible at the bottom left of the editor.

Press *Create*

Go back to the command line tab at <https://api.crc.testing:6443>

Type 'oc get configmaps'

Type 'oc describe configmap nodenewsconfig'

Go back to UI, select *Workloads/Deployment Configs*

Click on the *nodenews dc*

Click on *Environment*

In 'All Values from existing config maps' select nodenewsconfig from the pulldown on the left and add nodenewsconfig as the prefix

Click *Save*

Click on *Workloads/Pods* - if you are quick enough you will see the nodenews Pods being redeployed

By default when you create a Deployment Config in OpenShift, as part of an Application or as a standalone object via YAML, the Deployment Config will have two distinct triggers for

automation. These make the Deployment redeploy when a: the image that the Deployment Config is based on changes in the registry or b: the defined configuration of the Deployment changes via a change to the DC object itself

These are default behaviours but can be overridden. They are designed to make sure that the current deployment exactly matches the Images and the definition of the deployment by default.

Click on one of the Pods for *nodenews*

In the Pod Details page click on *Terminal*

In the terminal type ‘env | grep nodenewsconfig’

Note that the env variables from the config map have been expressed within the Pod as env variables (with the nodenewsconfig prefixed)

Go back to the command line tab at <https://api.crc.testing:6443> and enter the following commands:

```
mkdir material  
cd material  
git clone https://github.com/utherp0/workshop4  
cd workshop4/attendee  
cat testfile.yaml
```

All objects in Kubernetes/OpenShift can be expressed as yaml and this is a basic configmap for a file. What is very nice about the API and its object oriented nature is that you can export any object that you have the RBAC rights to view and import them directly back into OpenShift, which is what we will do now with the following commands in the terminal

```
oc create -f testfile.yaml  
oc get configmaps
```

Go back to the UI, select *Workloads/Deployment Configs*

Select *nodenews dc*

Click on *YAML*

In order to add the config-map as a volume we need to change the container specification within the deployment config.

Find the setting for ‘imagePullPolicy’. Put the cursor to the end of the line. Hit return. Underneath enter:

```
volumeMounts:  
  - name: workshop-testfile  
    mountPath: /workshop/config
```

Make sure the indentation is the same as for the ‘imagePullPolicy’.

Now in the ‘spec:’ portion we need to add our config-map as a volume.

Find ‘restartPolicy’. Put the cursor to the end of the line and press return. Underneath enter:

```
volumes:  
  - name: workshop-testfile  
configMap:  
  name: testfile  
  defaultMode: 420
```

Save the deployment config.

Click on *Workloads/Pods*. Watch the new versions of the nodenews application deploy.

When they finish deploying click on one of the Pods. Click on *Terminal*.

In the terminal type:

```
cd /workshop  
ls  
cd config
```

Note that we have a new file called ‘app.conf’ in this directory. This file is NOT part of the image that generated the container.

In the terminal type:

```
cat app.conf
```

This is the value from the configmap object expressed as a file into the running container.

In the terminal type:

```
vi app.conf
```

Press ‘i’ to insert, then type anything. Then press ESC. Then type ‘:wq’

You will not be able to save it. The file expressed into the Container from the configmap is ALWAYS readonly which ensures any information provided via the config map is controlled and immutable.

Type ':q!' to quit out of the editor

Dependency Injection of sensitive information using Secrets

The config map to be written as a file is actually written to the Container Hosts as a file, and then expressed into the running Container as a symbolic link. This is good but can be seen as somewhat insecure because the file is stored *as-is* on the Container Hosts, where the Containers are executed

For secure information, such as passwords, connection strings and the like, OpenShift has the concept of *Secrets*. These act like config maps *but* importantly the contents of the secrets are encrypted at creation, encrypted at storage when written to the Container Hosts and then unencrypted only when expressed into the Container, meaning only the running Container can see the value of the secret.

In the UI select *Workloads/Secrets*

Click on *Create*

Choose 'Key/Value Secret'

For 'Secret Name' give 'nodenewssecret'

Set 'Key' to 'password'

Set 'Value' textbox to 'mypassword'

Click 'Create'

When created click on the 'YAML' box in the Secrets/Secret Details overview

Note that the type is 'Opaque' and the data is encrypted

Click on 'Add Secret To Workload'

In the 'Select a workload' pulldown select the nodenews DC

Ensure the 'Add Secret As' is set to Environment Variables

Add the Prefix 'secret'

Click 'Save'

Watch the Pods update on the subsequent ‘DC Nodenews’ overview

When they have completed click on ‘Pods’

Choose one of the nodenews running pods, click on it, choose Terminal

In the terminal type ‘env | grep secret’

Understanding the Deployment Strategies

Click on *Workloads/Deployment Configs*

Click on the DC for *nodenews*

Scale the Application up to four copies using the up arrow next to the Pod count indicator

Once the count has gone to 4 and all the Pods are indicated as healthy (the colour of the Pod ring is blue for all Pods) select Action/Start Rollout.

Ignore the 404 (race condition) - click Back on the browser to go to the deployment config again

Watch the colours of the Pods

Deployments can have one of two strategies. This example uses the *Rolling* strategy which is designed for zero downtime deployments. It works but spinning up a single copy of the new Pod, and when that Pod reports as being healthy only then is one of the old Pods removed. This ensures that at all times the required number of replicas are running healthy with no downtime for the Application itself.

Click on *Workloads/Deployment Configs*, click on the nodenews DC again

Click on *Actions/Edit Deployment Config*

Scroll the editor down to the ‘spec:’ tag as shown below

```
spec:  
  strategy:  
    type: Rolling  
    rollingParams:
```

Change the type: tag of the strategy to Recreate as shown below

```
spec:  
  strategy:  
    type: Recreate
```

Click on Save

Click on *Workloads/Deployment Configs*, select nodenews dc

Click on ‘Action/Start Rollout’

Click back on the browser to go to the Deployment Configs/nodenews overview

Watch the colour of the Pod ring

In the case of a Recreate strategy the system ensures that NO copies of the old deployment are running simultaneously with the new ones. It deletes all the running Pods, irregardless of the required number of replicas, and when all Pods report as being fully deleted it will start spinning up the new copies. This is for a scenario when you must NOT have any users interacting with the old Application once the new one is deployed, such as a security flaw in the old Application

Application Deployment Configurations and DevOps Approaches

Introduction

DevOps is a much overloaded term that is used to describe a variety of concepts in the creation of modern applications. In a simple definition DevOps is concerned with the interface between development practices used for the creative elements of software engineering and the procedural rigour of operationally running applications in production. Clearly these concepts have different objectives so it is important for teams (whether tasked with development or operations) to have a good understanding of the concepts, objectives and concerns of their counterparts on the other side of the fence.

In terms of this workshop it is important to highlight the capabilities of containerised platforms with respect to the roll out of new versions of applications and how they get introduced to production and to end users. Terms such as blue/green deployments, black and white deployments, A/B deployments and canary deployments are used in various text books on *DevOps* principles and some of these areas will be used in this chapter of the workshop.

The activities in this workshop will introduce a new version of a simple application to use in two different manners.

Workshop Content

The application to be used in this workshop is a simple NodeJS REST interface. Version 1 exists on the master branch of the GIT Repository and version 2 exists on the experimental branch.

To begin the creation of the application use the following commands in an command line window. :

Creation of version 1

```
oc new-project master-slaveX
```

(Where X is your user number)

```
oc new-app --name slave-app-v1-0 https://github.com/marrober/slave-node-app.git
```

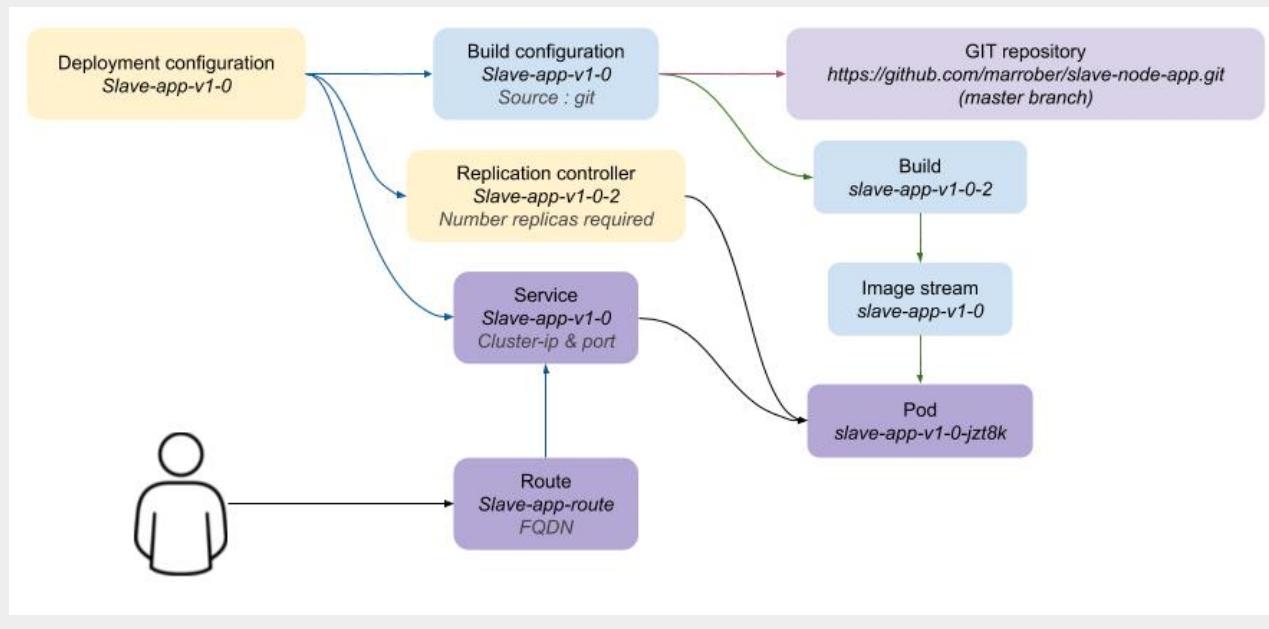
```
oc expose service slave-app-v1-0 --name="slave-app-route"
```

What got created ?

It is important to understand the artefacts and content that are created by the above commands within OpenShift:

- A new project is created called master-slaveX.
- A deployment configuration is created called slave-app-v1-0.
- A replication controller is created associated with the application and the built objects. This defines the scaling of the application in terms of the number of pods and the deployment strategy (rolling or recreate).
- A build configuration is created to compile the application source code which is extracted from the associated GIT repository.
- The build configuration will create a build instance which includes logs of the build process and a reference to the built image.
- The newly created image is deployed as a running pod (or pods depending on the number of pods defined in the deployment configuration).
- A service is created as part of the application creation process. The service endpoints will point to the pod(s) created by the build process. In the case of a rebuild and redeploy operation the service endpoints are updated to point to the new pods in accordance with the deployment strategy of the deployment configuration.
- The service is then exposed external to the cluster by the creation of a route. The route has a fully qualified domain name for access from machines outside the cluster.

The diagram below shows the above and how they relate together.



To identify the URL of the route execute the command shown below:

```
oc get route slave-app-route -o jsonpath='{"http://"}{.spec.host}{"/ip\n"}'
```

This will display a formatted URL with the *http://* part at the beginning which will be similar to :

<http://slave-app-route-master-slave.apps.cluster-xxxx-yyyy.xxxx-yyyy.example.opentlc.com/ip>

To test the application use the command line window to issue a curl command to:

```
curl -k <url from the above command>
```

The response should be as shown below (example ip address) :

"10.131.0.114 v1.0"

Creation of version 2

The development team now wants to introduce an experimental version 2 of the application and introduce it to the users in a number of different ways. The first action is to create the new build process for the experimental version using the command below.

```
oc new-app --name slave-app-v2-0 https://github.com/marrober/slave-node-app.git#experimental
```

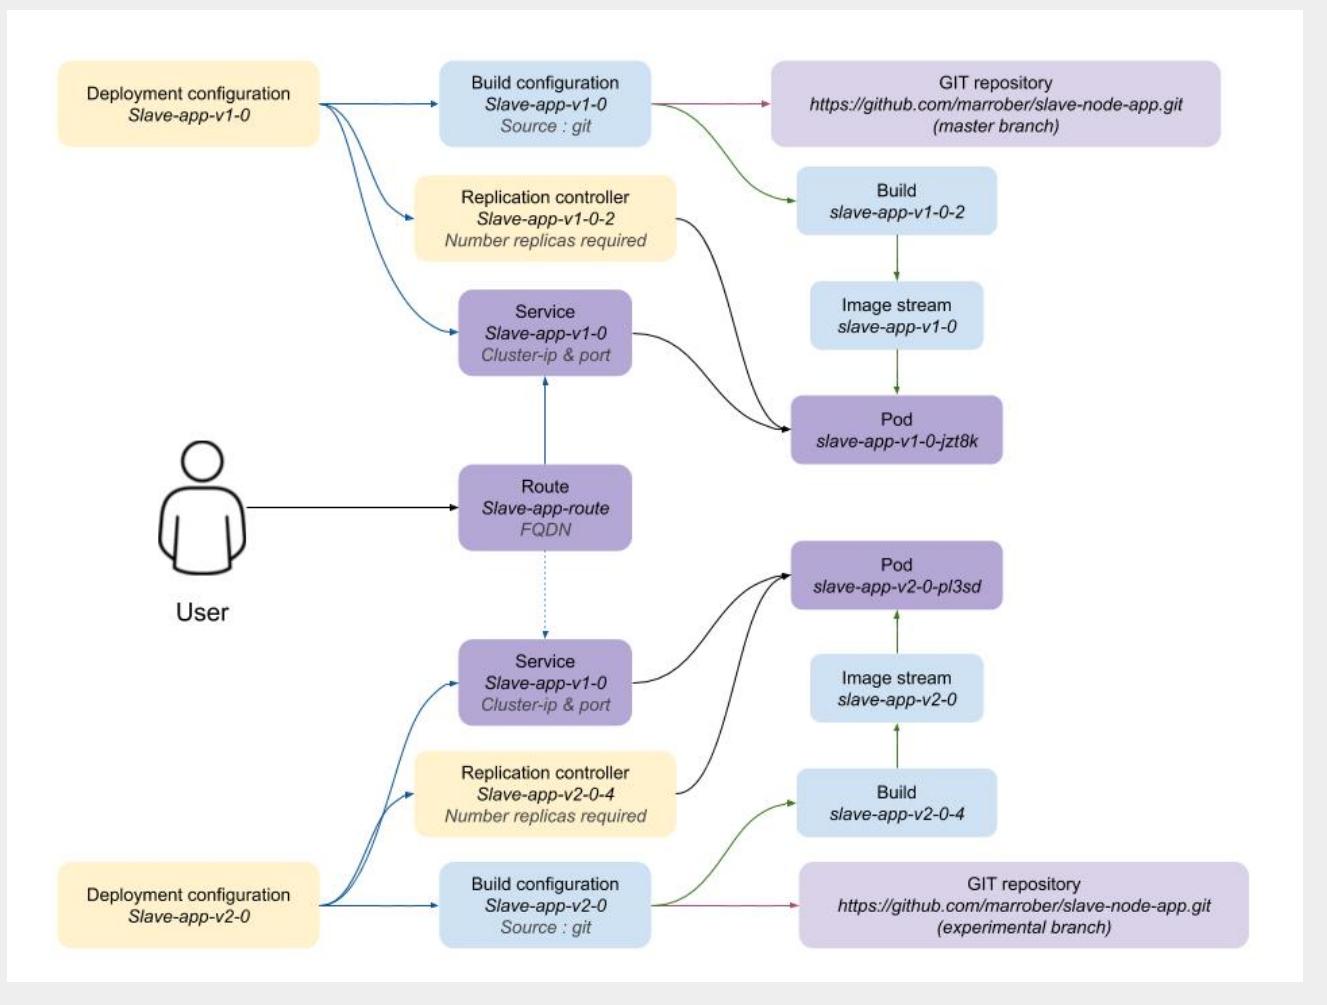
Note the use of the #experimental branch identifier in the end of the GIT repository. This syntax cannot be used through the web interface to openshift so any requirements for non-standard GIT connectivity must be done through the command line interface. It is also possible to configure the source-to-image capability to reference a specific sub directory of the repository using the --context option.

What got added to the project ?

New content was added to the project as a result of the above command specific to the experimental (v2) version of the application.

- A deployment configuration
- A replication controller
- A build configuration
- A running container in a pod
- A service

The diagram below shows the above and how they relate together.



At this point the version 2 application is operational but not accessible externally to the cluster.

Blue / Green Deployment

The benefit of creating the new version of the application alongside the old is that it is quick and

easy to migrate users to a new version of the application. It also allows teams to validate that the pods are running correctly.

Switching the route from v1 to v2 involves patching the route to change configuration. Before executing this operation open a new command window and execute the command below to send requests to the route every second. The responses received should all include the ip address of the pod and the version (v1) of the application.

Get the URI of the route using the command :

```
oc get route slave-app-route -o jsonpath='{http://}{.spec.host}{"/ip\n"}'
```

Copy the result of the above command and past it into the shell script below:

```
for i in {1..1000}; do curl -k <URI> ; echo ""; sleep 1; done
```

A series of reports of ip address and version 1 of the application will then start to scroll up the screen. Leave this running.

Switch back to the original command window and execute the command below to patch the route to version 2 of the application.

```
oc patch route/slave-app-route -p '{"spec":{"to":{"name":"slave-app-v2-0"}}}'
```

Switch back to the command window with the shell script running and you should see the responses have a new ip address and now report v2 of the application. This has completed a migration from the old version of the application to the new.

The details of the route patched by the above command are displayed by the command:

```
oc get route/slave-app-route -o yaml
```

The output of the above command is shown below, and the nested information from spec → to → name is easy to see.

```

apiVersion: route.openshift.io/v1
kind: Route
metadata:
  annotations:
    openshift.io/host.generated: "true"
  creationTimestamp: 2019-12-04T17:16:37Z
  labels:
    app: slave-app-v1-0
    name: slave-app-route
    namespace: master-slave
    resourceVersion: "884652"
  selfLink: /apis/route.openshift.io/v1/namespaces/master-slave/routes/slave-app-route
  uid: d4910fef-16b9-11ea-a6c5-0a580a800048
spec:
  host: slave-app-route-master-slave.apps.cluster-telf-c8e6.telf-c8e6.example.opentlc.com
  port:
    targetPort: 8080-tcp
  subdomain: ""
  to:
    kind: Service
    name: slave-app-v2-0
    weight: 100
  wildcardPolicy: None
status:
  ingress:
    - conditions:
        - lastTransitionTime: 2019-12-04T17:16:38Z
          status: "True"
          type: Admitted
      host: slave-app-route-master-slave.apps.cluster-telf-c8e6.telf-c8e6.example.opentlc.com
      routerCanonicalHostname: apps.cluster-telf-c8e6.telf-c8e6.example.opentlc.com
      routerName: default
      wildcardPolicy: None

```

Before moving to the A/B deployment strategy switch back to version v1 with the command:

```
oc patch route/slave-app-route -p '{"spec":{"to":{"name":"slave-app-v1-0"}}}'
```

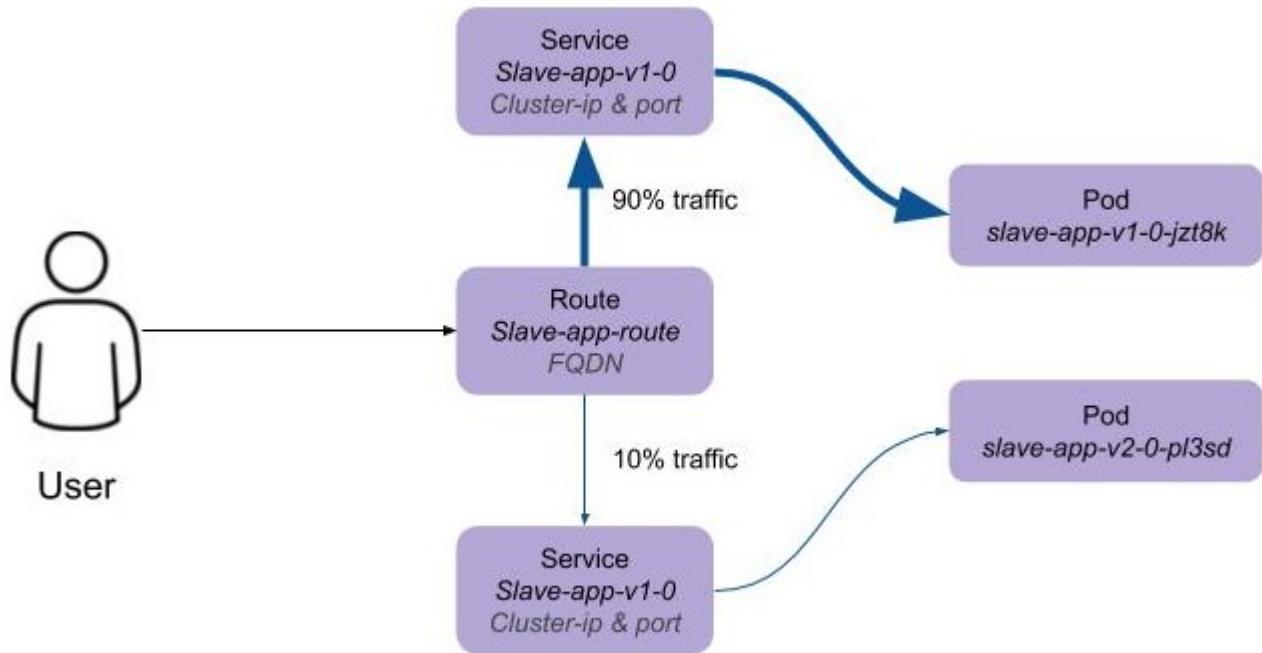
Confirm this has worked in the command window executing the shell script.

A/B Deployment

The benefit of an A/B deployment strategy is that it is possible to gradually migrate workload to

the new version. This example presents a simple process of gradually migrating a higher and higher percentage of traffic to the new version, however more advanced options are available for migrating traffic based on headers or source ip address to name just two. Red Hat OpenShift Service Mesh is another topic that is worth investigation if advanced traffic routing operations are required.

Gradually migrating traffic from v1 to v2 involves patching the route to change configuration as shown below.



To migrate 10% of traffic to version 2 execute the following command:

```
oc set route-backends slave-app-route slave-app-v1-0=90 slave-app-v2-0=10
```

Switch back to the command window running the shell script and after a short wait you will see the occasional report from version 2.

To balance the workload between the two versions execute the following command:

```
oc set route-backends slave-app-route slave-app-v1-0=50 slave-app-v2-0=50
```

Switch back to the command window running the shell script and after a short wait you will see a more even distribution of calls between versions 1 and 2.

The details of the route patched by the above command are displayed by the command:

```
oc get route/slave-app-route -o yaml
```

A section of the output of the above command is included below, showing the split of traffic between versions 1 and 2.

```
spec:  
  alternateBackends:  
    - kind: Service  
      name: slave-app-v2-0  
      weight: 50  
    host: slave-app-route-master-slave.apps.cluster-telf-c8e6.telf-c8e6.example.opentlc.com  
    port:  
      targetPort: 8080-tcp  
    subdomain: ""  
    to:  
      kind: Service  
      name: slave-app-v1-0  
      weight: 50
```

When satisfied that version 2 is working as required the following command will switch all traffic to that version and will remove the references to version 1 from the route.

```
oc set route-backends slave-app-route slave-app-v1-0=0 slave-app-v2-0=100
```

Cleaning up

From the OpenShift browser window click on *Advanced* and then *Projects* on the left hand side menu.

In the triple dot menu next to your own project (master-slaveX) select ‘Delete Project’ Type ‘master-slaveX’ (where X is your user number) such that the Delete button turns red and is active.

Press Delete to remove the project.

Understanding the Software Defined Network within OpenShift

Introduction

This chapter provides a developer-centric view of the fundamentals and capabilities of the SDN (Software Defined Network) that is used within OpenShift for Application connectivity.

If you are not already logged on go to the UI URL at <https://console-openshift-console.apps-crc.testing/> and logon as *userx* (x is the number provided to you) and password *openshift*.

The basics of Service Addressing

In this section we will create a couple of simple applications and show how they are represented via the Service and load-balancing of Pods within the SDN itself.

Open another tab in the browser and go to the terminal URL <https://api.crc.testing:6443> and again logon as *userx* with password *openshift*

Ensure you are on the Administrator View (top level, select Administrator)

Select Home/Projects

Click on ‘Create Project’

For the project name give it networktestx where x is your user number

Display Name and Description are optional

Click on *Create*

Using the top left selection switch the UI from Administrator to Developer

In the Topology Tab click on ‘From Catalog’

In the search box enter ‘node’. Select the Node.js option

Click ‘Create Application’

Default the image to version 10

For the git repo enter ‘<https://github.com/utherp0/nodenews>’

Set the *Application Name* to ‘application1-app’

Set the *Name* to *application1*

Click *Create*

Click on '+Add'

Click on 'From Catalog'

In the search box enter 'node'. Again select the Node.js option

Click 'Create Application'

For the git repo enter '<https://github.com/utherp0/nodenews>'

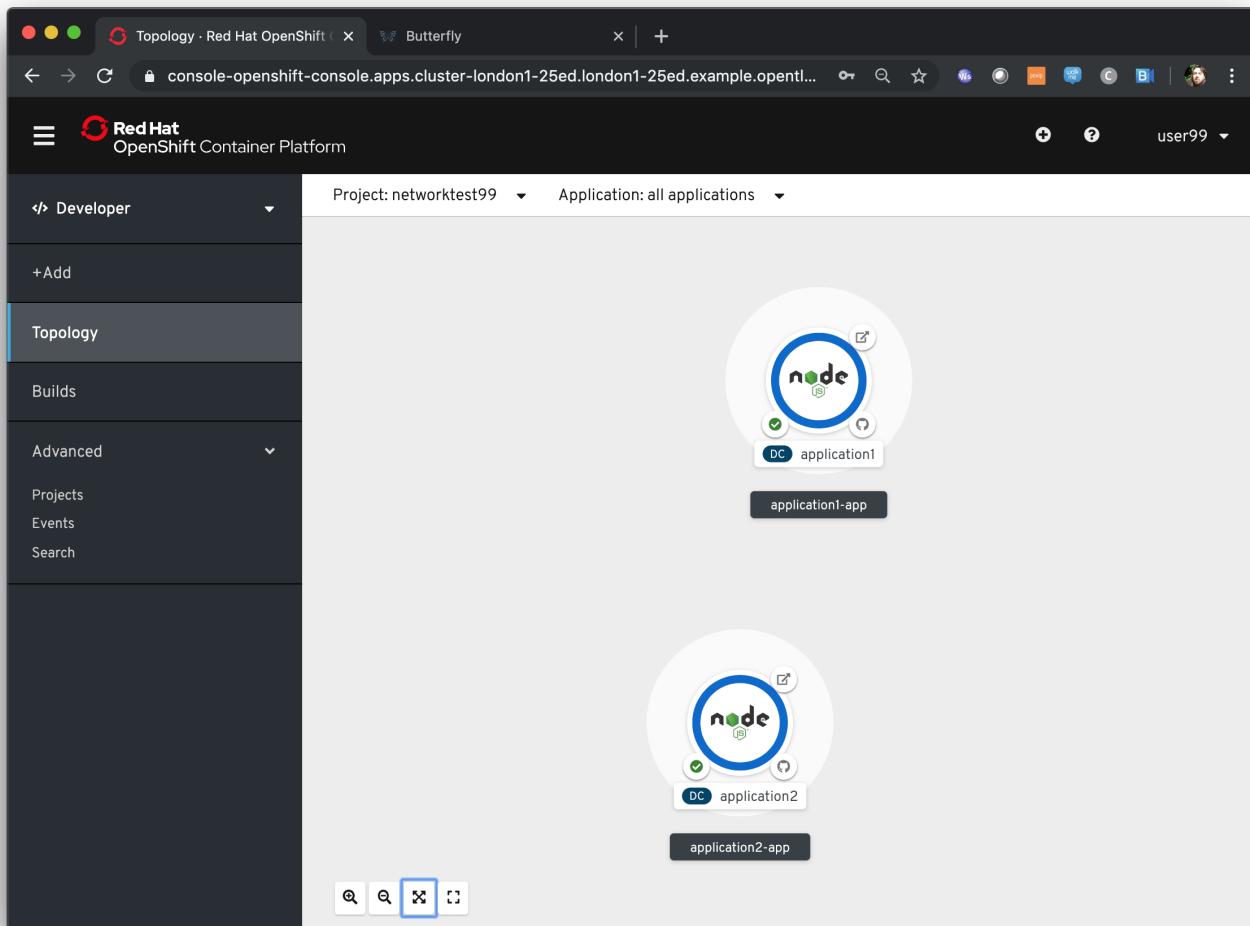
Choose *Create Application* from the Application pulldown

Set the *Application Name* to *application2-app*

Set the Name to 'application2'

Click *Create*

Wait until both applications have built and deployed. The Topology view will look similar to the image below:



Once the applications have created change the mode from Developer to Administrator using the top left selection point

Click on *Projects*. Select the networktestx project

Click on *Networking/Services*

The screenshot shows the Red Hat OpenShift Container Platform Services UI. The left sidebar is dark-themed and includes sections for Home, Projects, Search, Explore, Events, Operators, Workloads, Networking (which is selected), Services, Routes, Ingresses, Network Policies, Storage, Builds, and Administration. The main content area is titled 'Project: networktest99' and shows a list of 'Services'. There is a 'Create Service' button and a 'Filter by name...' input field. Two services are listed:

Name	Namespace	Labels	Pod Selector	Location
application1	networktest99	app=application1 app.kubernetes.io/c... = applica... app.kubernetes.io/i... = applica... app.kubernetes.io/n... = nodejs app.kubernetes.i... = applicatio... app.openshift.io/r... =nodejs app.openshift.io/r... =1...	app=application1, deploymentconfig=application1	172.30.54.191:8080
application2	networktest99	app=application2 app.kubernetes.io/c... = applica... app.kubernetes.io/i... = applica... app.kubernetes.io/n... = nodejs app.kubernetes.i... = applicatio... app.openshift.io/r... =nodejs app.openshift.io/r... =1...	app=application2, deploymentconfig=application2	172.30.196.252:8080

The Services UI shows the services currently available in the Project. Note that there is a Service per application, so in this case there are two Services, application1 and application2. The Location is a singular IP address for each Service within the SDN - more on this later.

Click on *Workloads/Deployment Configs*

Click on the application1 dc

Using the arrows next to the Pod indicator scale the application to three replicas

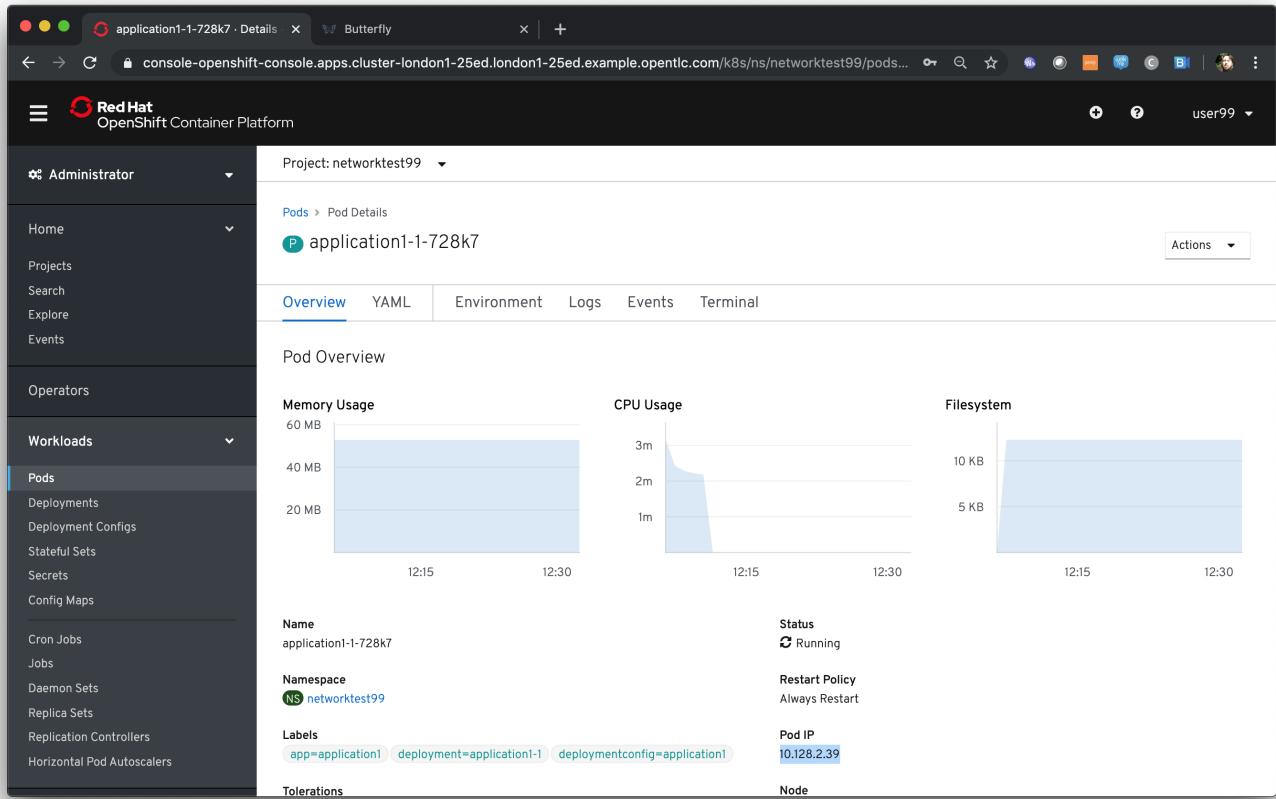
Click on 'Pods'

The screenshot shows the Red Hat OpenShift Container Platform interface. The left sidebar is titled "Administrator" and includes sections for Home, Projects, Search, Explore, Events, Operators, and Workloads. Under Workloads, the "Deployment Configs" option is selected. The main content area shows the "Deployment Configs > Deployment Config Details" for "application1". The "Pods" tab is active, showing three pods: "application1-1-728k7", "application1-1-8zprm", and "application1-1-hkdh4". All three pods are listed as "Running" with a "Ready" status. A filter bar at the top right allows filtering by name.

Name	Namespace	Pod Labels	Node	Status	Readiness
application1-1-728k7	networktest99	app=application1 deploy...=application1 deploym...=appli...	ip-10-0-137-221.eu-central-1.compute.internal	Running	Ready
application1-1-8zprm	networktest99	app=application1 deploy...=application1 deploym...=appli...	ip-10-0-137-221.eu-central-1.compute.internal	Running	Ready
application1-1-hkdh4	networktest99	app=application1 deploy...=application1 deploym...=appli...	ip-10-0-146-36.eu-central-1.compute.internal	Running	Ready

Click on the first active Pod for application1 listed

Scroll down and take a note of the IP address for the Pod



Click on *Networking/Services*

Click on the application1 service

Note that the IP address for the Service has NOT changed. Scaling the Application has no impact on the singular IP address for the Service, which in actuality acts as a load-balancer across ALL of the IP addresses for the Pods of that Application.

Using the shorthand Service name directly

Click on *Workloads/Pods*

Click on the first running Pod for application1 (the name will be application1-1-(something))

Click on Terminal

In the Terminal window type:

```
curl http://localhost:8080
```

You will see the webpage for the Application as a return from the curl statement. The Container sees itself as localhost. Also note that because we are calling from within the Container we use

the port address - if you were using the Routes their would be a Route for every Service, with no port address.

In the Terminal window type:

```
curl http://application1:8080
```

This is where it starts to get interesting. The Service *name* itself is exported into the network namespace of the Container so it can refer to it as a short name. The SDN translates the service name into the service IP - in reality this Container could be getting a webpage back from any of the Application Pods that satisfy this Service.

Using the Fully Qualified Domain Name for accessing Services

In the Terminal window type (and replace x with your number):

```
curl http://application1.networktestx.svc.cluster.local:8080
```

And this is the fully qualified version of the Service. by including the namespace/project name we can reach, effectively, any service on the SDN assuming the SDN has been configured to allow that. In this case we are just targeting our own Service from the application Container, now we will try the other application in the namespace.

In the Terminal window hit the up arrow to get the last command, edit the name and change application1 to application2, hit return at the end of the statement

You should a webpage. This is the output of the other application. This long format makes it easy to refer to other applications without having to leave and come back into the SDN (via a Route).

In the terminal type:

```
curl http://application2:8080
```

We can also connect to any of the Services hosted within the namespace/project by default

Ask the person sat next to you what their project name is and make a note of it

In the terminal type:

```
curl http://application1.(the project name from the person next to you).svc.cluster.local:8080
```

OpenShift Container Platform can be installed with two different modes of SDN. The first is subnet, which exposes all Services in all Namespace/Projects to each other. This instance has a subnet SDN which is why you should be able to call other peoples Services directly from your own via the internal FQDN address.

Controlling Access through Network Policies

Click on *Network/Network Policies*

For each of the policies listed click on the triple dot icon on the far right and choose ‘Delete Network Policy’.

Name	Namespace	Pod Selector
NP allow-from-all-namespaces	NS networktest99	All pods within networktest99
NP allow-from-ingress-namespace	NS networktest99	All pods within networktest99

The Network Policy tab should display ‘No Network Policies Found’.

Go to Workloads/Pods, click on one of the application1 Pods, choose Terminal

Repeat the ‘curl’ command listed above for the person sat next to you, i.e. curl their application1

Ensure you get a webpage

Go to Network/Network Policies

Click on ‘Create Network Policy’

Enter the following - remember to change YOURNUMBERHERE to your user number

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: example
  namespace: networktestYOURNUMBERHERE
spec:
  podSelector:
    matchLabels:
      app: application1
  ingress: []
```

Click ‘Create’

Wait until the person next to you has done the same

Click on *Workloads/Pods*, click on one of the application1 Pods, choose Terminal

Repeat the ‘curl’ command listed above for the person sat next to you, i.e. curl their application1

The call will eventually fail - feel free to hit Ctrl-C to interrupt

The creation of a Network Policy that prohibits ingress to the Application Service has stopped access to the Service from external namespaces AND internal Services.

Click on *Workloads/Pods*

Click on the active pod for application2

Click on Terminal

Type:

```
curl http://application1:8080
```

The call will eventually fail

This shows that the Service is prohibited even from Services in its own namespace/project. This application of Network Policy allows for fine-grain control of traffic egress/ingress at the Service level. The other installation mode for SDN for OpenShift 4 is with Network Policies enabled, with default Network Policies providing a fully multitenanted environment.

Click on *Projects*

On the triple dot icon on the far right for networktestxx select ‘Delete Project’

In the pop-up enter the name of the project (‘networktestxx’ with your number) and hit Delete

The RBAC model for Developers

Introduction

This chapter will introduce the attendee to the concepts around the Role Based Access Control model integrated into OpenShift. This will cover the use of Users, Service Accounts and Permissions.



For the majority of this chapter the attendee will need to pair up with another. If the numbers are odd the presenter should fill the gap of the missing attendee - pair the attendees up before the chapter starts and get them to exchange numbers - remember the numbers as MYUSER and PARTNERUSER

If you are not already logged on go to the UI URL <https://console-openshift-console.apps-crc.testing/> and logon as userx (x is the number provided to you) and password openshift.

Open another tab in the browser and go to the terminal URL <https://api.crc.testing:6443> and again logon as userx with password openshift

Ensure you are on the Administrator View (top level, select Administrator)

In the UI click on *Home/Projects*

Click on *Create Project*

For the project name give it rbactestx where x is your user number

Display Name and Description are optional

Click on *Create*

Examining Service Accounts

Once the project has been created, click on *Administration/Service Accounts*

Project: rbactest99

Name	Namespace	Secrets	Age
SA builder	rbactest99	2	a few seconds ago
SA default	rbactest99	2	a few seconds ago
SA deployer	rbactest99	2	a few seconds ago

Service Accounts are effectively virtual users - even though you have logged on when you do anything within the namespace, such as a build or a deployment, the namespace uses a virtual user, with appropriate permissions, to do the actual tasks internally, like pushing a built image to the registry, or pulling an image to start a Container.

Each project gets a Service Account by default, named default or to give it its full name system:serviceaccount:rbactestx:default (where x is your number). There are also Service Accounts for doing the appropriate actions, such as Builder and Deployer

In addition further Service Accounts can be created by the project owner and given additional security details and roles, such as being able to execute Containers as root, or giving Containers specific controlled sec-comp profiles (like changing the Group ID for the process)

Switch to the Developer view using the top left pulldown

In the Topology Tab click on 'From Catalog'

In the search box enter 'node'. Select the Node.js option

Click 'Create Application'

Default the image to version 10

For the git repo enter 'https://github.com/utherp0/nodenews'

Set the Application Name to ‘rbac’

Set the Name to ‘rbactest’

Click *Create*

Click on the Node pod indicator so the informational panel appears

Watch the build and ensure it completes

Once the build is complete ensure the deployment occurs. The Pod ring will turn dark blue

Switch back to Administrator view using the top left selection point

Click on *Home/Projects*

Click on Administration/Role Bindings

Note that we have one Role Binding shown at the moment, which is for Admin and applies to a kind of *user* and a subject name of *userx*. This is the binding of your user to the administration role within the project.

Click on the *System Role Bindings*

The screenshot shows the Red Hat OpenShift Container Platform web interface. The left sidebar is collapsed. The main header says "Role Bindings - Red Hat OpenS". The top right shows a user icon for "user99". The main content area has a title "Role Bindings" and a "Create Binding" button. A search bar says "Filter by role or subject...". Below is a table with four items:

Name	Role Ref	Subject Kind	Subject Name	Namespace	...
admin	admin	User	user99	rbactest99	...
system:deployers	system:deployer	ServiceAccount	deployer	rbactest99	...
system:image-builders	system:image-builder	ServiceAccount	builder	rbactest99	...
system:image-pullers	system:image-puller	Group	system:serviceaccounts:rbactest99	rbactest99	...

Note that with the addition of System Role Bindings to the screen we can now see the three Service Accounts also created. Deployer is part of the system:deployers binding, builder is part of the system:image-builders binding and the group of users under system:serviceaccounts:rbactestx (where x is your number) have the system:image-pullers binding.

Click on the (CR) admin Role Ref

This screen displays ALL the actions, via API, that this role has access to grouped by the API groups. This mapping is what controls what the user can do via the bindings they have.

Adding Role Bindings to your namespace/project

Click on *Administration/Role Bindings*

Click on *Create Binding*

Set Name to 'partneraccess'

Select 'rbactestx' from the Namespace pulldown; it should be the only one.

As a user with the admin role within the namespace/project you can add other users with role bindings within your project.

Select ‘Admin’ in the Role Name pulldown

Ensure ‘User’ is selected in the radiobox for Subject

Enter ‘userPARTNERUSER’ for the Subject Name

What we are doing is adding the user you have chosen to pair with as an admin role binding within your project.

Click *Create*

Ensure the partner has done the same with your userx

Click on Home/Projects

If the partner user has set the role binding appropriately you will now see two projects - your own and the other person’s

Click on the partner’s project (rbactestPARTNERUSER)

Change to the Developer view using the top left selection point

Ensure you can see the Topology page

Change back to the Administrator view using the top left selection point

Select *Workloads/Deployment Configs*

Ensure that the ‘rbactest’ DC shown has a Namespace that is the Partner’s project

Click on the DC rbactest

Using the arrows scale the deployment to 4 pods

Click on *Home/Projects* and select your project (rbactestMYUSER)

Click on ‘Role Bindings’ in the project overview pane

On the triple dot for ‘partneraccess’ choose ‘Delete’

Confirm deletion in the pop-up message box

Giving Users lower levels of permission

Click on *Administration/Role Bindings*

Click on *Create Binding*

Set Name to ‘partneraccess’

Choose the ‘rbactestMYUSER’ in the Namespace pulldown

Select ‘view’ in the Role Name pull down

Ensure the Subject radiobox is set to ‘User’

In the Subject Name enter the user name for the partner (userPARTNERUSER)

Click Create

Ensure the partner has done the same with your userx

Click on *Home/Projects*

Select the partner project (rbactestPARTNERUSER)

In the Project overview pane click on Role Bindings

You now do not have the appropriate access rights to interact with the role bindings as you only have View access to the target project,

Click on *Workloads/Deployment Config*

Click on the rbactest (DC)

Try and scale down the Pod to one pod

View access allows you to see the state of objects but NOT to change them.

Click on *Home/Projects*

In the triple dot menu next to the rbactestPARTNERUSER select ‘Delete Project’

Type ‘rbactestPARTNERUSER’ in the message box and press ‘Delete’

Note that you cannot delete the other persons project.

Hit Cancel

In the triple dot menu next to your own project (rbactestMYUSER) select ‘Delete Project’

Type ‘rbactestMYUSER’ in the message box and press ‘Delete’

Understanding Persistent Volumes

Introduction

If you are not already logged on go to the UI URL <https://console-openshift-console.apps-crc.testing/> and logon as userx (x is the number provided to you) and password openshift.

Open another tab in the browser and go to the terminal URL <https://api.crc.testing:6443> and again logon as userx with password openshift

Ensure you are on the Administrator View (top level, select Administrator)

Click on 'Create Project'

Enter 'pvtestx' for the name, where x is your user number

Enter 'PV Test' for the Display Name and Description

Hit 'Create'

Switch to the Developer UI (click on Developer in the top left pulldown)

Select 'From Catalog'

Search for 'node'

Click on the 'node.js' option

Click 'Create Application'

Leave the Builder Image Version as '10'

Enter the following for the Git Repo URL - '<https://github.com/utherp0/workshop4>'

Click on 'Show Advanced Git Options'

In Context Dir put '/nodeatomic'

In Application Name put 'nodeatomic'

In Name put 'nodeatomic'

Click on 'Create'

When the Topology page appears click on the node Pod to see the information window

The screenshot shows the Red Hat OpenShift Container Platform interface. The left sidebar has sections for Developer, Topology, Builds, Advanced, Projects, Events, and Search. The main area shows a topology diagram with a central node labeled 'node' and two surrounding nodes labeled 'nodeatomic'. Below the diagram is a search bar with four icons. The right panel is titled 'nodeatomic' and shows tabs for Overview and Resources (which is selected). Under Overview, there is a section for Pods with one entry: 'nodeatomic-1-b6qp4' (Running) with a 'View Logs' button. Under Builds, it shows a build named 'nodeatomic' with status 'Build #1 is complete (a minute ago)' and a 'View Logs' button. Under Services, it shows a service named 'nodeatomic' with 'Service port: 8080-tcp → Pod Port: 8080'. Under Routes, it shows a route named 'nodeatomic' with 'Location:'.

Ensure the build completes correctly

In the informational panel choose ‘Action/Edit Count’

Up the count to 4 and hit ‘Save’

Watch the Topology and ensure four copies of the Pod appear and are healthy

Adding a Persistent Volume to an Application

Switch to the Administrator view (top left pulldown)

Click on *Storage/Persistent Volume Claims*

Click on ‘Create Persistent Volume Claim’

Leave the Storage Class as ‘gp2’

Storage Classes are objects configured by the Administrator of the Cluster which allow you to request externalised persistent storage for your Applications. In this case we are using the default storage class for the Cluster which happens to be AWS storage

Set the Persistent Volume Claim Name to ‘testpvx’ where x is your user number

Ensure the Access Mode is set to ‘RWO’

The Access Mode for a Persistent Volume defines how the storage is offered to the Applications. If you choose RWO, Read/Write/Once, a single piece of storage is allocated which is shared across ALL instances of the Application. If you choose RWM, Read/Write/Many, the volume will be created independently on all Nodes where the Application runs, meaning the storage will only be shared by co-resident Pods on Nodes.

Set the size to 1Gi

Ensure the Use Label Selectors checkbox is not set

Click on ‘Create’

Click on Storage/Persistent Volume Claims

The status of the claim will sit at *Pending*. This is because the Persistent Volume itself is not created until it is assigned to a Deployment Config.

Click on *Workloads/Deployment Configs*

Click on the nodeatomic DC

Click on *Actions/Add Storage*

In the Add Storage page ensure ‘Use existing claim’ is checked

In the ‘Select Claim’ pulldown select the created claim (testpvx based on your user number)

Set the Mount Path to ‘/pvtest/files’

What this will do is to export the Persistent Volume into the file space of the Containers at the mount point stated.

Click on ‘Save’

As we have changed the configuration of the Deployment, and the default triggers are set to redeploy if the Image changes OR the configuration of the Deployment changes, a redeployment will now occur.

When the deployment has completed, click on ‘Pods’

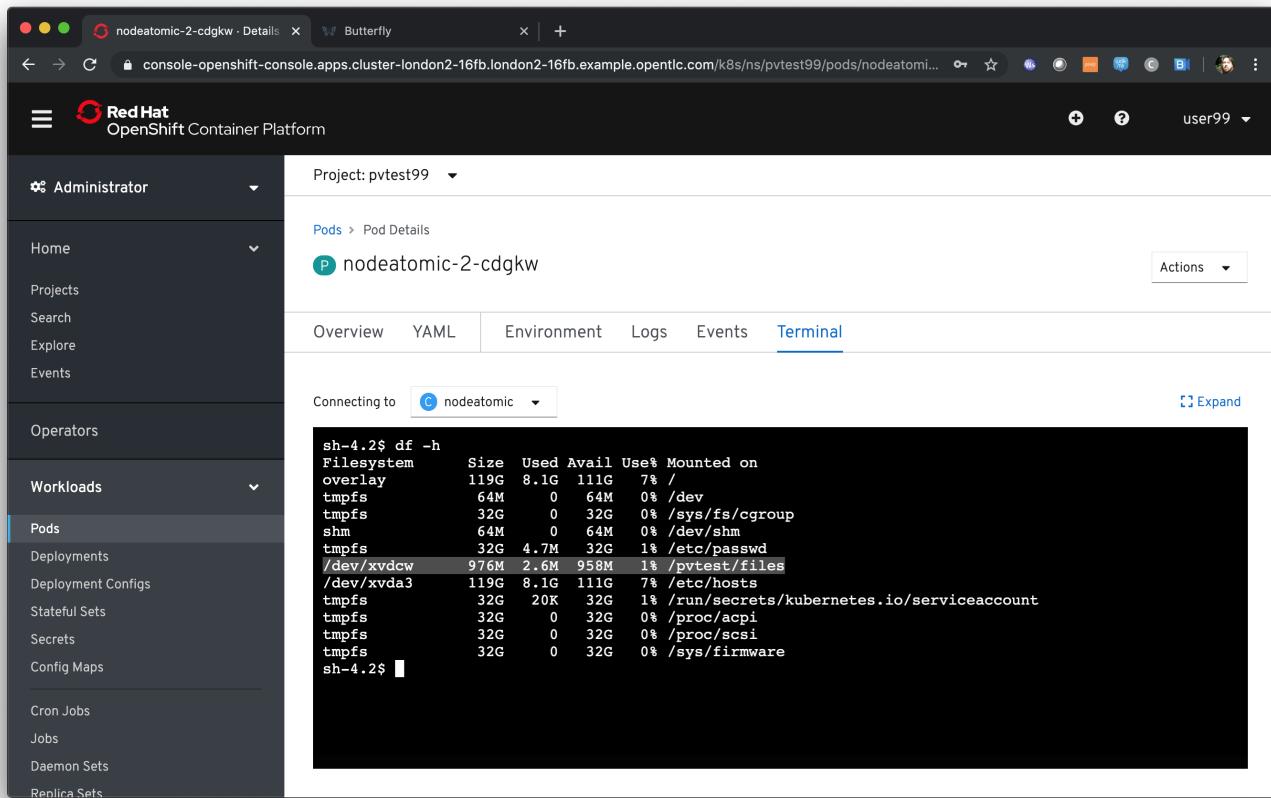
Select the first running Pod - take a note of its name (nodeatomic-2-xxxxx) - click on the Pod

Click on Terminal

In the terminal window type:

```
df -h'
```

*Note the *



In the terminal window type:

```
ls -alZ /pvtest/files
```

Click on *Networking/Routes*

On the Location of the nodeatomic RT (route) press the right mouse button (two fingers on the Mac mousepad) and select 'Open link in new tab'

Ensure the OpenShift NodeAtomic Example webpage is displayed

Add '/containerip' to the end of the URL in the browser window and hit return

Take a note of the address returned

Switch back to the OCP UI and choose Workloads/Pods

Click on **each** of the Pods until you find the one that has the IP returned by the webpage, take a note of

the Pod name (*1)

Go back to the tab with the nodeatomic webpage in it

Remove '/containerip' from the end of the URL and replace it with '/fileappend?file=/pvtest/files/webfile1.txt&text=Hello%20World' and then press return

Ensure the webservice returns 'Updated /pvtest/files/webfile1.txt with Hello World'

Switch back to the browser tab with the OCP UI in it. Select *Workloads/Pods* and click on the Pod with the name that matches the IP discovered in (*1)

Click on *Terminal*

In the terminal type:

```
cat /pvtest/files/webfile1.txt
```

Ensure 'Hello World' is displayed

The Webservice endpoint provided appends the given text to the given file.

Click on *Workloads/Pods*

Select another Pod (**NOT** the one that matched the IP from the (*1) step

Click on *Terminal*

In the terminal type:

```
cat /pvtest/files/webfile1.txt
```

Note that this separate Pod has the SAME file with the same contents

Switch back to the nodeatomic webservice browser tab

Alter the end of the URL to read 'Hello%20Again' and press return

Return to the OCP UI tab window (the terminal should still be active) and type:

```
cat /pvtest/files/webfile1.txt
```

Again note the file has been updated by another container but this container shares the same file system.

Close the web service browser tab

Demonstrating survivability of removal of all Pods

Click on *Workloads/Deployment Configs*

Click on the nodeatomic DC

Scale to ZERO pods

Ensure the Pod graphic displays zero running Pods.

Scale the deployment back up to ONE Pod using the arrows

When the Pod indicator goes to dark blue indicating the Pod has started, click on Pods

Select the one active Pod and click on it

Click on *Terminal*

In the terminal window type:

```
---
```

```
cat /pvtest/files/webfile1.txt
```

```
--
```

Note that the contents of the file have survived the destruction of ALL Pods

Click on *Home/Projects*

On the triple dot next to the ‘pvtestx’ project (where x is your user number) select Delete Project

In the pop-up type ‘pvtestx’ (where x is your user number) and hit Delete

Pod Health Probes

Introduction

liveness and readiness probes are Kubernetes capabilities that enable teams to make their containerised applications more reliable and robust. However, if used inappropriately they can result in none of the intended benefits, and can actually make a microservice based application unstable.

The purpose of each probe is quite simple and is described well in the OpenShift documentation here. The use of each probe is best understood by examining the action that will take place if the probe is activated.

Liveness : Under what circumstances is it appropriate to restart the pod?

Readiness : under what circumstances should we take the pod out of the list of service endpoints so that it no longer responds to requests?

Coupled with the action of the probe is the type of test that can be performed within the pod :

Http GET request : For success, a request to a specific http endpoint must result in a response between 200 and 399.

Execute a command : For success, the execution of a command within the container must result in a return code of 0.

TCP socket check : For success, a specific TCP socket must be successfully opened on the container.

Creating the project and application

Log on to cluster as userx, password openshift

Ensure you are on the Administrator View (top level, select Administrator)

The Administrator view provides you with an extended functionality interface that allows you to deep dive into the objects available to your user. The Developer view is an opinionated interface designed to ease the use of the system for developers. This workshop will have you swapping between the contexts for different tasks.

Click on *Create Project*

Name - ‘probesX’ where X is your assigned user number

Display Name - *Probes*

Description - *Liveness and readiness probes*

as shown in the image below:

The screenshot shows a 'Create Project' dialog box. It has three input fields: 'Name *' containing 'probes1', 'Display Name' containing 'Probes', and 'Description' containing 'Liveness and readiness probes.' At the bottom right are two buttons: 'Cancel' and a larger 'Create' button.

Create Project

Name *

probes1

Display Name

Probes

Description

Liveness and readiness probes.

Cancel Create

By default when you create a Project within OpenShift your user is given administration rights. This allows the user to create any objects that they have rights to create and to change the security and access settings for the project itself, i.e. add users as Administrators, Edit Access, Read access or disable other user's abilities to even see the project and the objects within.

In the top left of the UI, where the label indicates the view mode, change the mode from Administrator to Developer and select the Topology view.

Select 'From Catalog'

Enter 'node' in the search box

The various catalogue items present different configurations of applications that can be created. For example the node selections include a simple node.js application or node.js with a database platform that is pre-integrated and ready to use within the application. For this workshop you will use a simple node.js application.

Select 'Node.js'

A wizard page will then pop up on the right hand side with details of exactly what will be created through the process.

Source-2-Image

One of the most exciting features of OpenShift from a developer's perspective is the concept of S2I, or **Source-2-Image** which provides a standardised way of taking a base image, containing, for example, the framework for running a node.js application, a source code repository, containing the code of the application that matches the framework provided in the base image, and constructing and delivering a composite Application image to the Registry. Simply put this enables a developer to create source code for an application and OpenShift will convert that to a running application within a container with minimal effort. The process that you will use below is the Source-2-Image capability.

Click on 'Create Application'

The wizard process has a number of options that the user may elect to use. These include:

- Selecting a specific version of Node to use
- Selecting a GIT repository and choosing to use code from a specific directory within the repository.

Select the default offering for the Builder Image Version

For the GIT repository use : <https://github.com/marrober/slave-node-app.git>

In a separate browser tab go to <https://github.com/marrober/slave-node-app.git>

You can see that the GIT repository at this location only has a small number of files specific to the application. There is no content specific to how the application should be built or deployed into a container.

Close the github tab

Back at the OCP4.2 user interface complete the following information in the section titmed *General*.

For the Application drop down menu select *Create Application*.

For the Application name enter *Slave-application*.

For the Name field enter *node-app-slave*

The application name field is used to group together multiple items under a single grouping within the topology view. The name field is used to identify the specific containerised application.

Ensure the 'Create a route to the application' checkbox is checked.

Click Create

The user interface will then change to display the Topology view. This is a pictorial

representation of the various items (applications, databases etc.) that have been created from the catalogue and added to an application grouping.

Multiple application groupings can exist within a single project.

Viewing the running application

Click on the Icon marked 'Node'

A side panel will appear with information on the Deployment Configuration that has just been created. The overview tab shows summary information and allows the user to scale the number of pods up and down. The resources tab shows information on the building process, the pods and the services and route to reach the application (if these have been created).

Wait for the Build to finish, the Pod to change from Container Creating to Running.

Pod Colour Scheme

The colour scheme of the pods is important and conveys information about the pod health and status. The following colours are used :

- Running - Dark blue
- Not Ready - Mid blue
- Warning - Orange
- Failed - Red
- Pending - light blue
- Succeeded - Green
- Terminating - Black
- Unknown - Purple

When the build has completed the right hand side panel will show something similar to the image below. Note that the route will be different to that which is shown below.

DC node-app-slave

Actions ▾

Overview Resources

Pods

P node-app-slave-1-tw69j	Running	View Logs
--------------------------	---------	---------------------------

Builds

BC node-app-slave	Start Build
✓ Build #1 is complete (5 minutes ago)	View Logs

Services

S node-app-slave	Service port: 8080-tcp → Pod Port: 8080
------------------	---

Routes

RT node-app-slave	Location: http://node-app-slave-probes2.apps.cluster-newc-ea2b.newc-ea2b.example.opentlc.com ↗
-------------------	--

Click on the Tick at the bottom left of the Pod. Note that this display can also be shown by clicking on the 'View Logs' section on the right hand side panel.

The build log will show information on the execution of the source-2-image process.

Click on the arrow on the top right corner of the Pod, or click on the route URL shown in the right hand side resource details window. The application window will launch in a new browser window and should display text as shown below:

Hello - this is the simple slave REST interface v1.0

Liveness Probe

A number of probes will be created to show the different behaviours. The first probe will be a liveness probe that will result in the restart of the pod.

Since this work will be done using the oc command line you need to switch the current oc command line to work with the new project using the command:

```
oc project probesX
```

(Where X is the number that you used when you created the project)

To create the probe use the OC command line interface to execute the following command.

```
oc set probe dc/node-app-slave --liveness --initial-delay-seconds=30 --failure-threshold=1 --period-seconds=10 --get-url=http://:8080/health
```

The above probe will create a new liveness probe with the characteristics:

- Become active after 30 seconds
- Initiated a reboot after 1 instance of a failure to respond
- Probe the application every 10 seconds *Note that ordinarily a gap of 10 seconds between probes would be considered very long, but we use this time delay within the workshop to allow time for observing the behaviour of the probe.*
- Use the URL /health on the application at port 8080. Note that there is no need to specify a URL for the application.

The command line response should be as shown below.

```
deploymentconfig.apps.openshift.io/node-app-slave probes updated
```

Review the liveness probe information by executing the command:

```
oc describe dc/node-app-slave
```

The output of this command will include the following section that highlights the new liveness probe

```

Pod Template:
  Labels:  app=node-app-slave
            deploymentconfig=node-app-slave
  Containers:
    node-app-slave:
      Image:      image-registry.openshift-image-registry.svc:5000/probes2/node-app-
slave@sha256:bf377...241
      Port:       8080/TCP
      Host Port:  0/TCP
      Liveness:   http-get http://:8080/health delay=30s timeout=1s period=10s
                   #success=1 #failure=1
      Environment: <none>
      Mounts:     <none>
      Volumes:    <none>

```

Alternatively to view the probe in a different format use the command below:

```
oc get dc/node-app-slave -o yaml
```

Part of the output will show:

```

livenessProbe:
  failureThreshold: 1
  httpGet:
    path: /health
    port: 8080
    scheme: HTTP
  initialDelaySeconds: 30
  periodSeconds: 10
  successThreshold: 1
  timeoutSeconds: 1

```

To view the above information graphically then use the following steps:

Select the Topology view of the application.

Click on the pod in the centre of the screen to display the information panel on the right hand side. From the action menu on the right hand side click **Edit Deployment Configuration** as shown in the image below.

DC node-app-slave

Overview	<u>YAML</u>	Pods	Environment	Events
----------	-------------	------	-------------	--------

```

64      ports:
65        - containerPort: 8080
66          protocol: TCP
67        resources: {}
68        livenessProbe:
69          httpGet:
70            path: /health
71            port: 8080
72            scheme: HTTP
73          initialDelaySeconds: 30
74          timeoutSeconds: 1
75          periodSeconds: 10
76          successThreshold: 1
77          failureThreshold: 1
78        terminationMessagePath: /dev/termination-log
79        terminationMessagePolicy: File
80        imagePullPolicy: Always
81        restartPolicy: Always
82        terminationGracePeriodSeconds: 30
83        dnsPolicy: ClusterFirst
84        securityContext: {}
85        schedulerName: default-scheduler
86      status:
87        observedGeneration: 3
88        details:
89          message: config change
90          causes:
91            - type: ConfigChange

```

SaveReloadCancel

On the Deployment Configuration page that is displayed ensure that the YAML tab is selected and scroll down to aroundline 68 to see the YAML definition for the liveness probe. It is also possible to edit the parameters of the probe from this screen if necessary.

In order to execute the probe it is necessary to simulate a pod failure that will stop the application from responding to the health check. A specific REST interface on the application has been created for this purpose called `/ignore`.

Activation of the Liveness Probe

To view the activity of the probe it is necessary to open two windows.

Select the Topology view of the application.

Click on the arrow on the top right hand corner of the node icon to open the application URL in a new browser tab.

Back on the OpenShift browser tab, Click on the pod to open the details window on the right hand side and then click on the pod link on the resources tab. This will display a multi-tab window with details of the pod, select the events tab.

Switch to the application tab and put /ip on the end of the url and hit return. This will display the ip address of the pod.

Change the url to have /health on the end and hit return. This will display the amount of time that the pod has been running.

Change the url to have /ignore on the end and hit return. Quickly move to the browser tab showing the pod events and watch for the probe event.

The pod will restart after 1 failed probe event which takes up to 10 seconds depending on where the schedule is between the probe cycles. The events for the pod on the details screen will be similar to that shown below.

Streaming events... Showing 7 events

P node-app-slave-2-bh922 Generated from kubelet on ip-10-0-136-74.eu-central-1.compute.internal
Pulling image "image-registry.openshift-image-registry.svc:5000/probes2/node-app-slave@sha256:bf37789cff32a7f5ebb41de09f493db1374868efd7b0476c03cd50ec821cb241"

P node-app-slave-2-bh922 Generated from kubelet on ip-10-0-136-74.eu-central-1.compute.internal
Successfully pulled image "image-registry.openshift-image-registry.svc:5000/probes2/node-app-slave@sha256:bf37789cff32a7f5ebb41de09f493db1374868efd7b0476c03cd50ec821cb241"

P node-app-slave-2-bh922 Generated from kubelet on ip-10-0-136-74.eu-central-1.compute.internal
Created container node-app-slave

P node-app-slave-2-bh922 Generated from kubelet on ip-10-0-136-74.eu-central-1.compute.internal
Started container node-app-slave

P node-app-slave-2-bh922 Generated from kubelet on ip-10-0-136-74.eu-central-1.compute.internal
Liveness probe failed: Get http://10.128.2.173:8080/health: net/http: request canceled (Client.Timeout exceeded while awaiting headers)

The events after the firing of the liveness probe are the re-pulling and starting of the container image in a new pod.

Switch to the application tab and put /health on the end of the url and hit return. This will display the amount of time that the new pod has been running, which will understandably be a small number.

In order to experiment with the readiness probe it is necessary to switch off the liveness probe. This can either be done by changing the deployment configuration YAML definition using the web interface or by executing the following command line:

```
oc set probe dc/node-app-slave --liveness --remove
```

Readiness Probe

To create the probe use the OC command line interface to execute the following command.

```
oc set probe dc/node-app-slave --readiness --initial-delay-seconds=30 --failure-threshold=3 --success-threshold=1 --period-seconds=5 --get-url=http://:8080/health
```

The above command will create a new readiness probe with the characteristics:

- Become active after 30 seconds
- Remove the pod from the service endpoint after 3 instances of a failure to respond
- Cancel the removal of the pod and add it back to the service endpoint after 1 successful response
- Probe the application every 5 seconds
- Use the URL /health on the application at port 8080. Note that there is no need to specify a URL for the application.

The command line response should be as shown below

```
deploymentconfig.apps.openshift.io/node-app-slave probes updated
```

Review the probe created using the commands above:

```
oc describe dc/node-app-slave
```

and

```
oc get dc/node-app-slave -o yaml
```

View the state of the pod within the endpoints using the command below:

```
oc get ep/node-app-slave -o yaml
```

The output of the above command will list the details of the service endpoint which will include information on the pod to which the health probe is attached as shown below.

```

apiVersion: v1
kind: Endpoints
metadata:
  annotations:
    endpoints.kubernetes.io/last-change-trigger-time: 2019-11-26T16:04:50Z
  creationTimestamp: 2019-11-26T09:37:12Z
  labels:
    app: node-app-slave
    app.kubernetes.io/component: node-app-slave
    app.kubernetes.io/instance: node-app-slave
    app.kubernetes.io/name: nodejs
    app.kubernetes.io/part-of: master-slave
    app.openshift.io/runtime: nodejs
    app.openshift.io/runtime-version: "10"
  name: node-app-slave
  namespace: probes1
  resourceVersion: "1172051"
  selfLink: /api/v1/namespaces/probes1/endpoints/node-app-slave
  uid: 534139aa-1030-11ea-af1c-024039909e8a
subsets:
- addresses:
  - ip: 10.128.2.145
    nodeName: ip-10-0-136-74.eu-central-1.compute.internal
    targetRef:
      kind: Pod
      name: node-app-slave-5-hwj89
      namespace: probes1
      resourceVersion: "1172049"
      uid: ad6cc0e5-1043-11ea-af1c-024039909e8a
  ports:
  - name: 8080-tcp
    port: 8080
    protocol: TCP

```

The lines of interest above are:

```

subsets:
- addresses:
  - ip: 10.128.2.145

```

This shows that the address is currently part of the endpoint (it will participate in servicing requests) prior to the readiness probe activation.

Activation of the Readiness Probe

Select the Topology view of the application.

Click on the arrow on the top right hand corner of the node icon to open the application URL in a new browser tab (unless you already have one open).

On the OpenShift browser tab, click on the pod to open the details window on the right hand side and then click on the pod link on the resources tab. This will display a multi-tab window with details of the pod, select the events tab.

Switch to the application tab and put /ip on the end of the url and hit return. This will display the ip address of the pod.

Change the url to have /health on the end and hit return. This will display the amount of time that the pod has been running.

Change the url to have /ignore on the end and hit return. Quickly move to the browser tab showing the pod events and watch for the probe event.

The pod events will show a screen similar to that which is shown below.

The screenshot shows the 'Pods > Pod Details' interface for a pod named 'node-app-slave-4-cfp9c'. The 'Events' tab is selected. A green circular icon with a play/pause symbol indicates 'Streaming events...'. Below it, a red-bordered box highlights an event for 'node-app-slave-4-cfp9c' from 'Generated from kubelet on ip-10-0-146-213.eu-central-1.compute.internal'. The event details show a 'probes' status with a green circle and the text 'Nov 28, 3:29 pm'. A note at the bottom of the box states 'Readiness probe failed: Get http://10.131.0.183:8080/health: net/http: request canceled (Client.Timeout exceeded while awaiting headers)'. The top right corner of the interface has an 'Actions' dropdown menu.

Note that you will see the count of the readiness events incrementing every 5 seconds.

You will also see that the events continue counting up since readiness probes do not stop firing just because the pod has been removed from the endpoint list. It is important that they continue to probe since the conditions may change and it may be appropriate to add the pod back into the endpoint list.

View the state of the pod within the endpoints using the command below:

```
oc get ep/node-app-slave -o yaml
```

The output of the above command will list the details of the service endpoint which will include information on the pod to which the health probe is attached as shown below.

```
subsets:
- notReadyAddresses:
  - ip: 10.128.2.145
```

The subset of the command output shown above indicates that the address is now listed as ‘not ready’ and is not currently part of the endpoint.

Under production use conditions for the application may change and the pod may recover from the inability to respond to the readiness probe. If this happens then it will be added back to the endpoint list.

To simulate this the Node application has a REST endpoint at /restore. Since the pod is currently not receiving communications from outside the cluster the call to the restore endpoint is done from within the pod command window.

Switch to the OpenShift browser window that was showing the pod events.

Note that you will see a large number of pod readiness probe failures while you were reading the notes.

In the OpenShift Console choose Administrator View, then Workloads/Pods. Click on the Pod that is active and in the Pod information page click on the Terminal option.

Within the Pod Terminal enter the command :

```
curl -k localhost:8080/restore
```

You should see a response similar to that shown below (with a different IP address):

```
"10.128.2.146 restore switch activated"sh-4.2$
```

Now go back to the Terminal tab where you enter *oc* commands

View the state of the pod within the endpoints using the command below:

```
oc get ep/node-app-slave -o yaml
```

You should see that the line of interest, previously shown above, has changed back to that shown below:

```
subsets:  
- addresses:  
  - ip: <ip address of the pod>
```

On the OpenShift browser page switch back to the events tab and you should see that the readiness probe failure count is no longer increasing.

Finally, switch to the application browser page and change the URL to end in /health. You should see that the application has been running for some time (compared to the liveness probe that showed a restart had taken place) and it should be responding successfully to the health probe.

Cleaning up

From the OpenShift browser window click on *Advanced* and then *Projects* on the left hand side menu.

In the triple dot menu next to your own project (ProbesX) select ‘Delete Project’ Type ‘ProbesX’ (where X is your user number) such that the Delete button turns red and is active.

Press Delete to remove the project.